

Date of publication xxxx 00, 0000, date of current version xxxx 00, 0000.

Digital Object Identifier 00.0000/ACCESS.2017.DOI

Enhancing Thyroid Disease Diagnosis with Machine Learning and Counterfactual Explainable AI

MOHAMMAD ZAHIDUL ALAM¹, RUKSHANDA RAHMAN², HASAN MAHMUD SOZIB³,
HALIMA AHMED⁴, ADIB HOSSAIN⁵, ANSEENA ANEES SABEENA⁶, AFIA FAIROOZ
TASNIM⁷, FAHAD AHMED⁸, MD IMRAN SARKAR⁹, TIMOTEI ISTVÁN ERDEI¹⁰

¹Department of Information Technology, Westcliff University, USA (e-mail: m.alam.154@westcliff.edu)

²Department of Computer Science, Westcliff University, USA (e-mail: r.rahman.562@westcliff.edu)

³Department of Electrical and Electronic Engineering, Ahsanullah University of Science and Technology, Bangladesh (e-mail: sozib2019@gmail.com)

⁴Department of Science in Public Health, California State University, Los Angeles (Cal State LA), USA (e-mail: hahmed19@calstatela.edu)

⁵Department of Business Analytics, Trine University, USA (e-mail: adibhossain998@gmail.com)

⁶Department of Business Administration, Westcliff University, USA (e-mail: anseenaaneessabeena@gmail.com)

⁷Department of Public Health, California State University Long Beach, USA (e-mail: afia.tasnim009@gmail.com)

⁸Department of Science in Engineering Management, Trine University, USA (e-mail: fahadusa147@gmail.com)

⁹Department of Information Technology, Westcliff University, USA (e-mail: m.sarkar.129@westcliff.edu)

¹⁰Department of Vehicles Engineering, University of Debrecen, Hungary (e-mail: (no email provided))

Corresponding author: Hasan Mahmud Sozib (e-mail: sozib2019@gmail.com).

This research is funded by the X.

ABSTRACT Thyroid disease affects millions worldwide and is often underdiagnosed due to complex symptom overlaps and reliance on laboratory interpretation. Machine learning (ML) offers promise for automated diagnosis, but model transparency remains a barrier to clinical adoption. In this study, we develop an interpretable ML framework for thyroid disease classification using a real-world dataset from the UCI repository. We evaluate five classifiers—including Random Forest, CatBoost, and Support Vector Classifier—under a rigorous nested cross-validation scheme to ensure unbiased performance estimation. Our best models achieve near-perfect accuracy (up to 99.7%) and AUC (1.000) on a reduced set of ten consensus physiological features, demonstrating that high performance can be maintained with parsimonious inputs. To enhance trust and clinical utility, we integrate dual explainable AI (XAI) techniques: SHAP for local feature attribution and counterfactual analysis for actionable "what-if" reasoning. These explanations align closely with known endocrine physiology (e.g., TSH and FTI as primary drivers), supporting their face validity. While limited by class imbalance and lack of external validation, our approach combines strong predictive power with transparent decision-making, offering a foundation for trustworthy clinical decision support systems in thyroid care.

INDEX TERMS Goldenhar Syndrome, Object Detection, Artificial intelligence, Computer vision, MambaYOLO, Genetic Disorders, Deep learning.

THYROID is a butterfly-shaped endocrine gland situated anterior to the trachea, just below the larynx. Through secretion of triiodothyronine (T3) and thyroxine (T4) [1]–[3], it regulates basal metabolic rate, thermogenesis, cardiovascular tone, and neurocognitive function. Disruption of this axis—whether from autoimmune injury, iatrogenic causes, or micronutrient imbalance—produces a spectrum of disorders, including hypothyroidism, hyperthyroidism, thyroiditis, goiter, nodules, and malignancy [4]–[6]. The clinical burden is substantial: thyroid dysfunction is

common, often insidious in onset, and associated with preventable morbidity when diagnosis or treatment is delayed [6]–[10]. The gross anatomy of the thyroid gland and an illustration of goiter due to Hashimoto's thyroiditis are shown in Fig. 1. Standard diagnostic pathways combine clinical history and examination with laboratory testing (TSH, free/total T4, T3, thyroid autoantibodies) and selective imaging (ultrasound, radionuclide uptake/scan) [11]–[15]. While effective, these workflows face persistent limitations: inter-laboratory variability and thresholding differences,

observer dependence for image interpretation, class imbalance in routine cohorts that favors missed minority phenotypes, and latency between specimen collection and decision-making [16]–[18]. As a result, real-time, objective decision support that is both accurate and interpretable remains a priority for front-line care. Recent advances in machine learning (ML) have delivered strong performance on thyroid classification tasks spanning structured labs, ultrasound, histopathology, and multimodal data streams [19]. However, three gaps recur. First, many studies restrict to binary or three-way labeling and overlook clinically actionable subclasses. Second, evaluation protocols sometimes conflate model selection and assessment, inflating performance estimates. Third—and most critically for deployment—explanations are often post-hoc, coarse, or non-actionable, limiting clinicians' ability to understand why a model predicts a given label and what minimal changes could alter it [20], [21].




FIGURE 1: Anatomy of the thyroid gland and enlargement (goiter) typical of Hashimoto's thyroiditis. The thyroid lies anterior to the trachea below the larynx; autoimmune inflammation may enlarge the gland and disrupt hormone production [22].

The remainder of this paper is organized as follows. Section I provides a brief clinical background on thyroid disorders and current diagnostic challenges. Section II presents the dataset, feature selection strategies, and machine learning methodology, including nested cross-validation and interpretability components. Section ?? reports experimental findings and model comparisons under full and reduced feature sets. Section ?? discusses the study's limitations and outlines directions for future research. Finally, Section ?? concludes the paper and outlines future research directions.

I. BACKGROUND ON THYROID DISORDERS

A. CLINICAL OVERVIEW OF THYROID DISORDERS

Hyperthyroidism occurs when the thyroid gland produces an excessive amount of T4 hormone, leading to an overactive thyroid [23], [24]. Common causes include Graves' disease, thyroid nodules, and thyroiditis. Symptoms include tachycardia, unintentional weight

loss, arrhythmia, palpitations, and heat intolerance. If left untreated, complications may involve brittle bones, cardiac arrhythmias, and thyrotoxic crisis. Hypothyroidism is an underactive thyroid condition resulting from inadequate hormone production [25]–[27]. Etiologies include autoimmune disease, thyroid surgery, radiation therapy, or medication side effects. Symptoms range from fatigue, cold intolerance, and weight gain to depression and memory impairment. Untreated hypothyroidism can progress to severe outcomes such as myxedema or cardiovascular disease. Goiter refers to an abnormal enlargement of the thyroid gland [28]–[30]. Causes include iodine deficiency, autoimmune thyroiditis, nodules, or malignancy. Patients may experience neck swelling, difficulty swallowing, or hoarseness.

Thyroiditis denotes inflammation of the thyroid, encompassing Hashimoto's thyroiditis, subacute thyroiditis, and postpartum thyroiditis [31]–[33]. Clinical presentation includes neck pain, swelling, and variable hormone release depending on disease stage. Nodules are solid or cystic lumps within the gland [34]–[36]. Though often asymptomatic, large nodules can cause compressive symptoms or progress to hyperthyroidism or malignancy. Thyroid cancer arises from malignant transformation of follicular or parafollicular cells [37], [38]. Types include papillary, follicular, medullary, and anaplastic carcinoma. Although early stages may be asymptomatic, progressive disease manifests with neck lumps, voice changes, and regional lymphadenopathy.

B. MACHINE LEARNING APPROACHES

For thyroid disease prediction, machine learning methods have been widely explored across recent studies. Early prediction and stratification into cancer, hyperthyroidism, or hypothyroidism can substantially improve treatment timeliness and outcomes. In [39], a machine learning pipeline was proposed to predict the presence of BRAF mutations in thyroid nodules using ultrasound from 96 cases; Random Forest (RF), Logistic Regression (LR), and Support Vector Machines (SVM) exceeded 60% accuracy. A related effort combined fine-needle aspiration (FNA) cytology with ultrasound features to reduce false negatives in thyroid cancer; RF outperformed gradient methods and decision trees. Building on ultrasound-derived predictors, LR and LASSO were applied in [40] to select malignancy-associated features, followed by RF with a scoring scheme, achieving >80% accuracy. In [41], ten classifiers were benchmarked; Extra Trees yielded 84% accuracy. SVM-based multi-class detection across four thyroid states was reported in [42] with 83.37% accuracy. Decision Trees (DT) were used in [43] to develop a tool (MLTDD) that achieved 98.7% accuracy. Broader molecule-screening pipelines leveraging LR, RF, SVM, Gradient Boosting Machines (GBM), and deep neural networks (DNN) to identify disruptors of thyroid homeostasis were explored in [44]. Feature-engineering comparisons across backward/forward/bidirectional selection and tree-based

selection were studied in [45], where Extra Trees–selected features paired with RF reached 99% accuracy and enabled predictions for non-thyroidal syndrome and autoimmune entities. Ensembling and treatment-guidance studies include: meta-classification with boosting/bagging/stacking/voting on a large, complex dataset with careful train–test splits and varied k values in [46]; a treatment-direction predictor for hypothyroid patients using SMOTE and Extra Trees achieving 84% across accuracy/precision/recall/F1 in [47]; and further RF-centered comparative studies reporting 94.8% in [48]. Using the UCI thyroid dataset, six ML algorithms with 10-fold CV were compared in [49]; MLP led on accuracy/recall/F1. A hybrid LDA–KNN–ANFIS pipeline was introduced in [50], combining dimensionality reduction, weighted preprocessing, and neuro-fuzzy inference. A multi-class framework with feature engineering and Extra Trees plus three DL baselines (batch 16, 100 epochs) was reported in [51], evaluated via accuracy, precision, recall, F1, 10-fold CV, SD, and confusion matrices. Additional empirical studies also report strong performance: random forest outperforming peers for hyperthyroidism identification in [52]; multi-model comparisons for three-class thyroidism in [53] with MLP best at 96.4%; and separate analyses distinguishing hyper- vs. hypothyroidism [54] and XGBoost superiority on a UCI dataset [55]. Complementary works explored fuzzy models [56], expert systems [57], swarm optimization for thyroid function analysis [58]. Table 1 provides the comparative analysis of the existing works.

TABLE 1: State-of-the-Art Models for Thyroid Disease Prediction

Ref	Year	Dataset	Model	Results
[?]	2021	UCI new-thyroid dataset	Logistic Regression, Decision Tree, k-Nearest Neighbor	kNN: 96.875% accuracy; DT: 87.5% accuracy; LR: 81.25% accuracy
[?]	2022	UCI thyroid disease dataset	Random Forest (with MLFS feature selection)	RF: 0.99 accuracy, 0.99 F1 score
[?]	2024	Clinical thyroid disease dataset (1232 samples)	Stacking Ensemble (LR, SVM, KNN, DT, ANN)	Stacking Ensemble: 99.9% accuracy, 99.9% sensitivity, 99.8% specificity
[?]	2024	UCI thyroid dataset	AdaBoost (with DE optimization)	AdaBoost: 0.998 accuracy
[?]	2023	UCI thyroid dataset	Random Forest (with LASSO)	RF: 99% accuracy, 97% precision, 92% recall, 95% F1 score
[?]	2024	UCI thyroid dataset	Ensemble ML (Hard Voting on RF and DT)	Ensemble: 99.71% accuracy, 100% sensitivity
[?]	2022	UCI thyroid dataset	Decision Tree, Random Forest	DT: 99.5758% accuracy; RF: 99.3107% accuracy

II. METHODOLOGY




FIGURE 2: Proposed pipeline for thyroid disease prediction: data cleaning, feature engineering, nested cross-validation for model training/selection, and downstream explainability (SHAP) and counterfactual analysis.

A. OUR PROPOSED APPROACH

As illustrated in Figure 2, our proposed pipeline begins with preprocessing (data cleaning and feature engineering), followed by nested cross-validation for model training, tuning, and evaluation, and concludes with post hoc explainability (SHAP) and counterfactual analysis. We treat each supervised learner as a distinct hypothesis family that offers a complementary inductive bias over the same clinical signal. Concretely, we train five models—Support Vector Classifier (SVC), Random Forest (RF), Extremely Randomized Trees (Extra Trees), AdaBoost, and CatBoost—on a single real dataset of thyroid cases with three outcome classes (Normal, Hypothyroid, Hyperthyroid). To probe the trade-off between parsimony and capacity, we evaluate two feature configurations: (i) *FULL* (all available predictors) and (ii) *TOP10* (the consensus set selected via multi-criteria feature selection). Each learner supplies its own “view” of the data: margin-based separation (SVC), bagged axis-aligned partitions (RF), high-randomness ensembles (Extra Trees), stage-wise additive boosting (AdaBoost), and ordered gradient boosting with symmetric trees (CatBoost). This diversity is intentional; it allows us to assess stability of conclusions across fundamentally different decision mechanisms [?], [?].

All modeling adheres to a strict nested cross-validation protocol (Algorithm 1). In the outer loop (stratified 5-fold), we obtain unbiased estimates of generalization. Within each outer training fold, a stratified inner loop (3-fold) performs hyperparameter selection via grid search with accuracy as the primary score. Model selection is thus confined to the inner loop, while the outer loop is used exclusively for evaluation. This design prevents information leakage from validation to test partitions, and yields robust performance summaries as means and standard deviations across outer folds [?], [?].

Algorithm 1 Nested Cross-Validation for Thyroid Classification (Single Real Dataset)

Require:

Single dataset $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$
 Feature sets $F \in \{\text{FULL}, \text{TOP10}\}$
 Models $M = \{\text{SVC}, \text{RF}, \text{ExtraTrees}, \text{AdaBoost}, \text{CatBoost}\}$
 Grids $\{G_m\}_{m \in M}$

Ensure:

Outer-CV mean metrics and best hyperparameters per model and feature set

1: Preprocessing:

Standardize continuous labs (e.g., TSH, FTI, T3, TT4, T4U) with training-fold statistics; keep binary indicators as $\{0, 1\}$.

No data leakage: fit scalers only on inner/train folds.

2: For each feature set $F \in \{\text{FULL}, \text{TOP10}\}$

3: Select columns X_F according to F ; set labels y (3 classes)

4: Define outer CV as stratified 5-fold with fixed seed

5: for each outer fold $t = 1, \dots, 5$ do

6: Split $(X_F, y) \rightarrow (\text{trainout}, \text{testout})$ stratified by class

7: for each model $m \in M$ do

8: **Model Pipeline:** Build pipeline Π_m

9: if $m = \text{SVC}$ then

10: // SVC requires scaling; probability=True for AUC

11: scaler $\rightarrow \text{SVC}(C, \text{kernel}, \text{probability=True})$

12: else

13: // Trees/boosting do not require scaling

14: passthrough $\rightarrow m$

15: end if

16: Define inner CV as stratified 3-fold on trainout

17: **Training:** Perform grid search over G_m using inner CV; scoring=accuracy; refit on full trainout to obtain $\hat{\theta}_m^{(t)}$

18: **Evaluation:** Predict on testout: get $\hat{y}_m^{(t)}$ and class probabilities $\hat{P}_m^{(t)}$

19: Compute metrics on testout:

20: Accuracy, weighted Precision/Recall/F1 (per scikit-learn definitions)

21: Multiclass AUC (OvR):

$$\text{AUC}_m^{(t)} = \frac{1}{K} \sum_{c=1}^K \text{AUC}_{-y=c} \left(\hat{P}_m(y = c | x) \right)$$

22: Store $\{\text{metrics}_m^{(t)}, \hat{\theta}_m^{(t)}\}$

23: end for

24: end for

25: **Final Aggregation:** For each model $m \in M$, aggregate over outer folds: report mean $\pm$ std of all metrics; summarize best hyperparameters (e.g., by fold-wise mode or best mean inner-CV score)

=0

B. DATASET DESCRIPTION

For this study, we utilized the **Thyroid Disease dataset** from the UCI Machine Learning Repository, a well-established benchmark for thyroid disorder classification. This dataset comprises anonymized patient records associated with various thyroid conditions, categorized into three class labels: Normal, Hypothyroid, and Hyperthyroid. The dataset includes a total of 7,200 patient samples, as outlined below:

- **Normal** ($N = 6,666$)
- **Hyperthyroid** ($N = 368$)
- **Hypothyroid** ($N = 166$)

Due to the absence of etiological markers (e.g., TRH or ACTH levels), the dataset does not differentiate between primary and secondary subtypes of hypothyroidism or hyperthyroidism. As a result, this study focuses on high-level thyroid status classification, deferring subtype differentiation to future work, which will incorporate more granular data, such as pituitary function tests.

The dataset consists of 21 clinical features, which include demographic indicators (e.g., age, sex), treatment history (e.g., on_thyroxine, thyroid_surgery), and laboratory biomarkers (e.g., TSH, T3, TT4, T4U, FTI). These features are commonly used in thyroid function diagnosis. The dataset is notably imbalanced, with most patients labeled as normal or hypothyroid. The features include both *binary clinical attributes* (e.g., on_thyroxine, sick, pregnant) and *continuous biochemical markers* (e.g., TSH, T3, TT4, T4U, FTI).

For data preparation, we merged the two structured text files provided by UCI—ann-train.data and ann-test.data—into a single dataset. The final dataset includes a total of 7,200 samples, each with 21 input features, and a multiclass target variable indicating thyroid disease classification.

TABLE 1: BASELINE CHARACTERISTICS BY THYROID CLASS

Table 2 presents a summary of baseline patient characteristics stratified by thyroid class. Key findings include:

- **Age** was similar across classes, with normalized means ranging between 0.518 and 0.554.
- **Sex** showed a statistically significant difference ($p < 0.001$), with a higher proportion of females in the hypothyroid group.
- **Biomarkers** such as TSH, T3, TT4, T4U, and FTI displayed expected trends and highly significant differences ($p < 0.001$).
- **Medication-related variables**, such as on_thyroxine and on_antithyroid_medication, also showed significant differences across classes.
- Sparse features like pregnant, lithium, goitre, and tumor were retained for completeness and future explainability analysis.

Continuous variables were scaled to the range of $[0, 1]$, and binary variables were encoded as “Yes”/“No” to enhance

clarity and interpretability in downstream modeling and XAI visualization.

The demographic and clinical characteristics of three groups (Group 1, $n = 166$; Group 2, $n = 368$; Group 3, $n = 6666$) and their combined total ($n = 7200$) were analyzed, with p-values assessing differences across groups. The mean age was similar across groups (0.536 ± 0.191 , 0.554 ± 0.188 , 0.518 ± 0.189 , respectively; overall 0.521 ± 0.189), with medians ranging from 0.540 to 0.590 and no significant difference ($p = 0$). Sex distribution showed a female predominance (79.5%, 78.0%, 68.9% for Groups 1, 2, and 3), with Group 3 having a higher male proportion (31.1%) compared to Groups 1 (20.5%) and 2 (22.0%; $p < 0.001$). Thyroid-related variables revealed significant differences in thyroxine use (10.2%, 0%, 13.8%; $p < 0.001$), driven by Group 2's absence of thyroxine, and in querying hypothyroidism (14.5%, 13.9%, 6.0%; $p < 0.001$), with Groups 1 and 2 showing higher rates. No significant differences were found for querying thyroxine ($p = 0.42$), antithyroid medication ($p = 0.72$), hyperthyroidism ($p = 0.76$), thyroid surgery ($p = 0.09$), I-131 treatment ($p = 0.25$), sickness ($p = 0.17$), pregnancy ($p = 0.09$), goitre ($p = 0.17$), or tumors ($p = 0.98$).

Psychiatric conditions were slightly more prevalent in Group 3 (5.1%) than Groups 1 (0.6%) and 2 (2.7%; $p = 0.01$). Thyroid function tests showed significant differences ($p < 0.001$) in TSH (means: 0.0963, 0.0137, 0.00210 mIU/L), T3 (0.00965, 0.0173, 0.0204), TT4 (0.0326, 0.0940, 0.112), T4U (0.103, 0.0978, 0.0977), and FTI (0.0323, 0.0969, 0.116), indicating distinct thyroid profiles, with Group 1 having higher TSH, suggesting possible hypothyroidism, and Group 3 showing higher T3 and TT4, potentially indicating hyperthyroidism.

C. EXPLORATORY DATA ANALYSIS

To gain insights into the distribution, correlation, and relationships between variables in the merged thyroid dataset, we conducted a comprehensive exploratory data analysis (EDA), summarized in this section.

1) Class Distribution

Figure 3 presents the distribution of disease classes within the dataset. The overall distribution highlights the predominance of hyperthyroid (normal thyroid function) cases, followed by a smaller but clinically important set of hypothyroid and hyperthyroid cases. The imbalance observed across classes is typical of real-world clinical data and underscores the importance of synthetic data generation and careful sampling to mitigate model bias. Specifically, the hypothyroid group accounts for a noticeable portion of the dataset and is presumed to include both primary hypothyroidism (characterized by insufficient thyroid hormone production due to gland dysfunction) and secondary hypothyroidism (typically resulting from pituitary or hypothalamic causes). These subtypes, while related, may differ in severity and clinical presentation,



FIGURE 3: Thyroid class distribution: Normal, Hyperthyroid, Hypothyroid

which can influence predictive modeling. However, the dataset does not explicitly distinguish these subtypes due to the absence of etiological markers (e.g., TRH or ACTH levels). In contrast, the hyperthyroid group is comparatively smaller in size, reflecting the lower prevalence of this disorder in the source population. This group is presumed to include subtypes such as primary hyperthyroidism (commonly due to Graves' disease or toxic nodular goiter) and secondary or drug-induced hyperthyroidism. Although less frequent, hyperthyroid cases are clinically significant because of their association with increased cardiovascular risk and metabolic complications. The dataset does not explicitly distinguish subtypes of both hypothyroidism and hyperthyroidism, which are unevenly represented. While primary hypothyroidism is presumed to dominate within the hypothyroid category, hyperthyroid subtypes are fewer but clinically diverse. This distribution provides both opportunities and challenges for classification tasks, especially for models that must capture nuanced patterns across underrepresented categories. We therefore model the three main classes (normal, hypothyroid, hyperthyroid), acknowledging that clinical subtypes exist but are not separable in this cohort, and defer subtype differentiation to future work with more granular data.

As illustrated in Figure 3, the dataset exhibits a notable class imbalance, with a majority of samples labeled as Normal, while the Hyperthyroid and Hypothyroid classes are relatively underrepresented. The dataset exhibits moderate class imbalance: Normal (92.6%), Hypothyroid (6.5%), Hyperthyroid (0.9%). While high-performing models achieved strong accuracy, this raises concern about potential bias toward majority classes. This skewed distribution raises the question of whether class balancing techniques—such as Synthetic Minority Over-sampling Technique (SMOTE)—are necessary to mitigate potential model bias toward the majority class. In the current phase of the analysis, we hypothesize that high-performing classifiers, particularly ensemble models such as Random Forest and CatBoost, may still yield robust results without the need

TABLE 2: Baseline Characteristics by Thyroid Class

Variable	Hypothyroid (N=166)	Hyperthyroid (N=368)	Normal(N=6666)	Overall (N=7200)	P-value
Age (years)					0
Mean (SD)	0.536 (0.191)	0.554 (0.188)	0.518 (0.189)	0.521 (0.189)	
Median [Min, Max]	0.560 [0.0100, 0.890]	0.590 [0.0100, 0.910]	0.540 [0.0100, 0.970]	0.550 [0.0100, 0.970]	
Sex					< 0.001
Female	132 (79.5%)	287 (78.0%)	4590 (68.9%)	5009 (69.6%)	
Male	34 (20.5%)	81 (22.0%)	2076 (31.1%)	2191 (30.4%)	
On thyroxine					< 0.001
No	149 (89.8%)	368 (100%)	5743 (86.2%)	6260 (86.9%)	
Yes	17 (10.2%)	0 (0%)	923 (13.8%)	940 (13.1%)	
Query thyroxine					0.42
No	166 (100%)	363 (98.6%)	6560 (98.4%)	7089 (98.5%)	
Yes	0 (0%)	5 (1.4%)	106 (1.6%)	111 (1.5%)	
On antithyroid medication					0.72
No	165 (99.4%)	365 (99.2%)	6578 (98.7%)	7108 (98.7%)	
Yes	1 (0.6%)	3 (0.8%)	88 (1.3%)	92 (1.3%)	
Sick					0.17
No	165 (99.4%)	352 (95.7%)	6407 (96.1%)	6924 (96.2%)	
Yes	1 (0.6%)	16 (4.3%)	259 (3.9%)	276 (3.8%)	
Pregnant					0.09
No	166 (100%)	368 (100%)	6588 (98.8%)	7122 (98.9%)	
Yes	0 (0%)	0 (0%)	78 (1.2%)	78 (1.1%)	
Thyroid surgery					0.09
No	162 (97.6%)	368 (100%)	6569 (98.5%)	7099 (98.6%)	
Yes	4 (2.4%)	0 (0%)	97 (1.5%)	101 (1.4%)	
I-131 treatment					0.25
No	160 (96.4%)	362 (98.4%)	6557 (98.4%)	7079 (98.3%)	
Yes	6 (3.6%)	6 (1.6%)	109 (1.6%)	121 (1.7%)	
Query hypothyroid					< 0.001
No	142 (85.5%)	317 (86.1%)	6269 (94.0%)	6728 (93.4%)	
Yes	24 (14.5%)	51 (13.9%)	397 (6.0%)	472 (6.6%)	
Query hyperthyroid					0.76
No	158 (95.2%)	344 (93.5%)	6203 (93.1%)	6705 (93.1%)	
Yes	8 (4.8%)	24 (6.5%)	463 (6.9%)	495 (6.9%)	
Lithium					0.49
No	166 (100%)	362 (98.4%)	6581 (98.7%)	7109 (98.7%)	
Yes	0 (0%)	6 (1.6%)	85 (1.3%)	91 (1.3%)	
Goitre					0.17
No	166 (100%)	368 (100%)	6607 (99.1%)	7141 (99.2%)	
Yes	0 (0%)	0 (0%)	59 (0.9%)	59 (0.8%)	
Tumor					0.98
No	162 (97.6%)	360 (97.8%)	6494 (97.4%)	7016 (97.4%)	
Yes	4 (2.4%)	8 (2.2%)	172 (2.6%)	184 (2.6%)	
Hypopituitary					1
No	166 (100%)	368 (100%)	6665 (100.0%)	7199 (100.0%)	
Yes	0 (0%)	0 (0%)	1 (0.0%)	1 (0.0%)	
Psych					0.01
No	165 (99.4%)	358 (97.3%)	6325 (94.9%)	6848 (95.1%)	
Yes	1 (0.6%)	10 (2.7%)	341 (5.1%)	352 (4.9%)	
TSH (mIU/L)					< 0.001
Mean (SD)	0.0963 (0.113)	0.0137 (0.0119)	0.00210 (0.00507)	0.00486 (0.0230)	
Median [Min, Max]	0.0530 [0.00620, 0.530]	0.00970 [0.00610, 0.143]	0.00150 [0, 0.177]	0.00170 [0, 0.530]	
T3					< 0.001
Mean (SD)	0.00965 (0.00592)	0.0173 (0.00543)	0.0204 (0.00746)	0.0200 (0.00754)	
Median [Min, Max]	0.00960 [0.000500, 0.0260]	0.0170 [0.000500, 0.0410]	0.0201 [0.000500, 0.180]	0.0201 [0.000500, 0.180]	
TT4					< 0.001
Mean (SD)	0.0326 (0.0213)	0.0940 (0.0215)	0.112 (0.0345)	0.109 (0.0360)	
Median [Min, Max]	0.0316 [0.00200, 0.0890]	0.0900 [0.0440, 0.150]	0.109 [0.00500, 0.600]	0.107 [0.00200, 0.600]	
T4U					< 0.001
Mean (SD)	0.103 (0.0178)	0.0978 (0.0166)	0.0977 (0.0192)	0.0978 (0.0191)	
Median [Min, Max]	0.103 [0.0610, 0.159]	0.0965 [0.0560, 0.165]	0.0960 [0.0170, 0.233]	0.0970 [0.0170, 0.233]	
FTI					< 0.001
Mean (SD)	0.0323 (0.0204)	0.0969 (0.0196)	0.116 (0.0345)	0.113 (0.0361)	
Median [Min, Max]	0.0325 [0.00200, 0.0640]	0.0950 [0.0649, 0.155]	0.112 [0.00415, 0.642]	0.111 [0.00200, 0.642]	

The figure displays five box plots, each representing a different thyroid lab value: TSH, T3, TT4, T4U, and T4I. The x-axis for all plots categorizes the data into three thyroid classes: Normal (green), Hypothyroid (orange), and Hyperthyroid (blue). The y-axis for each plot represents the concentration of the respective lab value. Individual data points are overlaid on the box plots to show the distribution and identify outliers.

- TSH Distribution by Thyroid Class:** The y-axis is labeled 'TSH' and ranges from 0.0 to 0.6. The Normal group has a median around 0.02, Hypothyroid around 0.01, and Hyperthyroid around 0.08. There are several outliers in the Normal and Hyperthyroid groups.
- T3 Distribution by Thyroid Class:** The y-axis is labeled 'T3' and ranges from 0.000 to 0.175. The Normal group has a median around 0.04, Hypothyroid around 0.03, and Hyperthyroid around 0.02. The Normal group shows a wide range of values with many outliers.
- TT4 Distribution by Thyroid Class:** The y-axis is labeled 'TT4' and ranges from 0.0 to 0.6. The Normal group has a median around 0.08, Hypothyroid around 0.05, and Hyperthyroid around 0.03. There are outliers in the Normal group.
- T4U Distribution by Thyroid Class:** The y-axis is labeled 'T4U' and ranges from 0.00 to 0.20. The Normal group has a median around 0.10, Hypothyroid around 0.10, and Hyperthyroid around 0.11. There are outliers in the Normal group.
- T4I Distribution by Thyroid Class:** The y-axis is labeled 'T4I' and ranges from 0.00 to 0.6. The Normal group has a median around 0.08, Hypothyroid around 0.05, and Hyperthyroid around 0.03. There are outliers in the Normal group.

Figure 10 displays 18 bar charts arranged in a 3x6 grid, showing the distribution of categorical variables across three thyroid classes: Normal (blue), Hypothyroid (orange), and Hyperthyroid (green). The variables are: sex, no_hypertension, overweight_moderation, peak, no_hypertension, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, 123_headrest, and 123_headrest. The y-axis represents the count of observations for each category.

in Hypothyroid cases, T3 and TT4 levels are elevated in Hyperthyroid cases, and T4U shows a broader distribution with higher medians in Hyperthyroid cases. These trends align with clinical expectations, where TSH is typically low in Hyperthyroid and high in Hypothyroid conditions, while T3, TT4, and T4U reflect thyroid hormone activity.

Figure 6 visualizes how categorical clinical indicators vary by thyroid class. Notably, features such as on thyroxine, on antithyroid medication, and query hypothyroid exhibit clear differences in their distribution across the classes. Most categorical features have a majority of 'No' (0) responses, but certain attributes become more prevalent in the diseased groups, aiding interpretability.

7

5) Feature Selection

To reduce dimensionality, enhance interpretability, and mitigate overfitting in our thyroid classification pipeline, we systematically evaluated a set of filter, wrapper, and embedded feature selection techniques. Given our dataset $D = \{(x_i, y_i)\}_1^N$ with $N = 19,998$ patient instances, where $x_i \in \mathbb{R}^{21}$ is the clinical predictor vector, and $y_i \in \{\text{Normal, Hypothyroid, Hyperthyroid}\}$, we define the goal of feature selection as identifying an optimal subset $S \subseteq \{1, \dots, 21\}$ that maximizes classification performance.

6) Mutual Information (MI)

Mutual information quantifies the reduction in uncertainty of the class label Y given a feature X_j . For a discrete or discretized feature X_j and target Y , the mutual information is defined as:

$$I(X_j; Y) = \sum_{x \in X_j} \sum_{y \in Y} p(x, y) \log \frac{p(x, y)}{p(x)p(y)} \quad (1)$$

In practice, we estimate $p(x, y)$ using empirical frequencies and apply the ‘mutual information function from scikit-learn’ to all $j = 1, \dots, 21$ features. The top k features with the highest $I(X_j; Y)$ are retained.

7) ANOVA F-Statistic (Numerical Features)

For continuous lab variables such as TSH, T3, TT4, and T4U, we employ the one-way ANOVA F-statistic to measure inter-class separability:

$$F_j = \frac{\sum_{k=1}^K n_k (\bar{x}_{k,j} - \bar{x}_j)^2}{\sum_{k=1}^K \sum_{i \in C_k} (x_{i,j} - \bar{x}_{k,j})^2} \quad (2)$$

where $K = 3$ (number of thyroid classes), n_k is the size of class C_k , $\bar{x}_{k,j}$ is the mean of feature j in class k , and $\bar{x}_j$ is the global mean. This method assumes normally distributed residuals and homoscedasticity.

8) Chi-Squared Test (Categorical Features)

For binary/categorical indicators such as sex, on thyroxine, tumor, etc., we compute the chi-squared statistic to test independence with thyroid class:

$$\chi^2 = \sum_r \sum_c \frac{(O_{rc} - E_{rc})^2}{E_{rc}} \quad (3)$$

where O_{rc} is the observed frequency of feature level r and class c , and E_{rc} is the expected frequency under independence. The expected frequency E_{rc} is calculated as:

$$E_{rc} = \frac{(\text{row}_r \text{ total}) \times (\text{col}_c \text{ total})}{N} \quad (4)$$

9) Recursive Feature Elimination (Wrapper-Based)

Let $f(x; \theta)$ be a base estimator (e.g., Logistic Regression), and let $L(\theta; D)$ be the cross-validated loss. Recursive Feature Elimination (RFE) operates by iteratively removing features that contribute least to performance:

$$S^{(t+1)} = S^{(t)} \setminus \{j^*\} = \arg \min \frac{\partial L}{\partial \theta} \quad (5)$$

We initialize with $S^{(0)} = \{1, \dots, 21\}$ and reduce iteratively until a stopping threshold $|S^{(t)}| = k$ is met.

10) L1-Regularized Logistic Regression (Embedded)

Using a multi-class L1-penalized logistic regression, we induce sparsity in the feature coefficients. The optimization objective is:

$$\min_{\theta} - \sum_{i=1}^N \sum_{c=1}^3 \mathbb{I}(y_i = c) \log \left(\frac{\exp(\theta_c^\top x_i)}{\sum_{c'=1}^3 \exp(\theta_{c'}^\top x_i)} \right) + \lambda \sum_{c=1}^3 \|\theta_c\|_1 \quad (6)$$

Features with zero-valued coefficients $\theta_{cj} = 0$ for all c are removed.

11) Random Forest Feature Importance (Gini Impurity)

In tree-based models, we compute feature importance based on the average reduction in Gini impurity across all trees and splits:

$$\text{Imp}(X_j) = \sum_{s \in \text{splits on } X_j} \Delta_i s \cdot n_s \quad (7)$$

where T is the number of trees, s denotes a split node, $\Delta_i s$ is the impurity decrease at node s , and n_s is the number of samples at node s .

The top 15-ranked features obtained from each of the six methods were compiled and saved for transparency and reproducibility. These full-feature ranking tables are presented in Appendix B, with the following organization:

- Mutual Information (See Table 13)
- ANOVA F-test (See Table 14)
- Chi-Squared Statistic (See Table 15)
- L1-Regularized Logistic Regression (See Table 16)
- Recursive Feature Elimination (RFE) (See Table 17)
- Random Forest Importance (See Table 18)

To enable visual comparison across methods, we additionally present a consolidated grid-based visualization of the top 15 features from each technique in Figure 7. This highlights both consensus and divergence across different algorithmic perspectives on feature relevance, informing subsequent model design decisions.

To identify the most consistently informative features across diverse selection criteria, we adopted a consensus-based strategy. Let $M = \{\text{MI, ANOVA, } \chi^2, \text{L1, RFE, RF}\}$ denote the set of six feature selection methods. For each method $m \in M$, we define the ranked top- k feature set as:

$$F(k) = \{x_j \in X : \text{Rank}_m(x_j) \leq k\} \quad (7)$$

where $\text{Rank}_m(x_j)$ is the importance rank of feature x_j under method m , and X is the full feature space.

We then define the consensus feature subset as the intersection across all methods:




FIGURE 7: Visual comparison of top 15 feature rankings across six feature selection methods: Mutual Information, ANOVA F-test, Chi-Squared, L1-Regularized Logistic Regression, RFE, and Random Forest.

$$F(k)_{\text{consensus}} = \bigcap_{m \in M} F(k)_m \quad (8)$$

In our case, we set $k = 15$ and extracted $F(15)_{\text{consensus}}$, selecting the top $d = 10$ most frequently recurring features across all $F(15)_m$ sets shown in Table 4. This approach emphasizes robustness and model-agnostic agreement in feature utility.

D. MODEL DEVELOPMENT

We selected these five models to represent diverse algorithmic paradigms and inductive biases:

- **SVC**: Margin-based linear separation, sensitive to scaling but robust under high dimensionality.
- **Random Forest & Extra Trees**: Bagged tree ensembles offering stability and native resistance to overfitting.
- **AdaBoost**: Sequential boosting sensitive to weak learners and noise.
- **CatBoost**: Gradient boosting optimized for categorical variables and reduced overfitting via ordered boosting.

This diversity allows us to evaluate performance consistency across fundamentally different decision mechanisms. Tree-based models were prioritized due to their compatibility with exact SHAP computation (TreeSHAP), which ensures higher fidelity explanations compared to approximation methods required for other architectures.

1) Problem Setup and Notation

Let $D = \{(x_i, y_i)\}_1^n$ denote the training dataset, where $x_i \in \mathbb{R}^{10}$ is the i -th input feature vector derived from clinical and demographic variables, and $y_i \in \{-1, +1\}$ is the corresponding binary label indicating the presence (+1) or absence (1) of the target thyroid disorder. The feature space is defined as:

$$x = [\text{TSH}, \text{FTI}, \text{T3}, \text{TT4}, \text{on thyroxine},$$

TABLE 4: Top 10 Consensus Features and Their Clinical Interpretations in Thyroid Classification

Feature	Description and Relevance to Thyroid Function
TSH	Thyroid Stimulating Hormone — Primary regulator of thyroid activity. Elevated TSH typically indicates hypothyroidism; suppressed TSH is seen in hyperthyroidism. Most discriminative marker for thyroid dysfunction.
FTI	Free Thyroxine Index — Adjusted estimate of free T4 levels, accounting for thyroid-binding proteins. Helps differentiate between euthyroid and thyrotoxic states. Useful when total T4 or T3 binding abnormalities are suspected.
T3	Triiodothyronine — Active thyroid hormone. Often elevated in hyperthyroidism. Sensitive in early thyrotoxicosis when T4 may still be normal.
TT4	Total Thyroxine — Total amount of T4 (bound + unbound). Indicates hormone production levels but is influenced by binding proteins. Helps distinguish overt vs. subclinical disease.
on thyroxine	indicates whether the patient is currently taking levothyroxine (synthetic T4). Important for detecting treated hypothyroidism or replacement therapy. Impacts the interpretation of lab values.
hypothyroid flag	Query hypothyroid flag if the patient is suspected to have hypothyroidism but not definitively diagnosed. Can help identify borderline or subclinical cases. Useful for early-stage pattern recognition.
psych	Psychiatric history flag — Thyroid dysfunction, especially hypothyroidism, can present with depressive symptoms, while hyperthyroidism may present with anxiety or agitation. Adds behavioral dimension.
thyroid surgery	indicates a history of thyroidectomy. Strongly affects hormonal regulation and may signal post-operative hypothyroidism or need for hormone replacement.
T4U	Thyroxine Uptake — Proxy for thyroid-binding capacity. Used with TT4 to calculate FTI. Low in hypothyroidism; high in hyperthyroidism. May reflect abnormal protein binding states.
sex	Biological sex — Thyroid diseases are more common in females (esp. autoimmune hypothyroidism). Can moderate clinical presentation and hormone reference ranges. [Ref: NIH Women's Health Initiative].

query hypothyroid, psych, thyroid surgery, T4U, sex^T, where continuous variables are standardized to zero mean and unit variance, and binary indicators (e.g., on thyroxine) are encoded as $\{0, 1\}$, with sex typically mapped as 0 (female) and 1 (male). All features are assumed to lie in a measurable space $X \subseteq \mathbb{R}^{10}$, and the label space is $Y = \{-1, +1\}$.

Our objective is to learn a classifier $f : X \rightarrow Y$ that minimizes the expected risk:

$$R(f) = \mathbb{E}_{(x,y) \sim P} [\ell(f(x), y)], \quad (9)$$

where $\ell : Y \times Y \rightarrow \mathbb{R}_{\geq 0}$ is a loss function, and P is the unknown joint distribution over $X \times Y$. Since P is inaccessible, we minimize the empirical risk:

$$\hat{R}_n(f) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i), \quad (10)$$

regularized appropriately depending on the model class.

We employ five state-of-the-art supervised learning algorithms: Support Vector Classifier (SVC), Random Forest (RF), Extra Trees, AdaBoost, and CatBoost. Each is formalized below with full mathematical rigor.

2) Support Vector Classifier (SVC)

The SVC seeks a maximum-margin hyperplane in a high-dimensional reproducing kernel Hilbert space (RKHS) $\mathcal{H}$ induced by a Mercer kernel $K : X \times X \rightarrow \mathbb{R}$. Let $\varphi : X \rightarrow \mathcal{H}$ be the associated feature map such that $K(x_i, x_j) = \langle \varphi(x_i), \varphi(x_j) \rangle_{\mathcal{H}}$.

The primal optimization problem is:

$$\min_{w \in \mathcal{H}, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \frac{1}{2} \|w\|_{\mathcal{H}}^2 + C \sum_{i=1}^n \xi_i, \quad (11)$$

subject to:

$$y_i (\langle w, \varphi(x_i) \rangle_{\mathcal{H}} + b) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad \forall i = 1, \dots, n,$$

where w is the normal vector to the hyperplane, b the bias, ξ the slack variables, and $C > 0$ the regularization parameter controlling the trade-off between margin maximization and classification error.

By the representer theorem, the solution admits the form:

$$w = \sum_{i=1}^n \alpha_i y_i \varphi(x_i), \quad \alpha_i \in \mathbb{R},$$

leading to the dual problem:

$$\max_{\alpha} \left(\sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j K(x_i, x_j) \right), \quad (12)$$

subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^n \alpha_i y_i = 0$. We employ the radial basis function (RBF) kernel:

$$K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|_2^2), \quad \gamma > 0. \quad (13)$$

The final decision function is:

$$f_{\text{SVC}}(x) = \text{sign} \left(\sum_{i \in S} \alpha_i y_i K(x_i, x) + b \right), \quad (14)$$

where $S = \{i : \alpha_i > 0\}$ is the set of support vectors.

3) Random Forest (RF)

Random Forest is an ensemble of T decorrelated decision trees $\{h_t(x; \Theta_t)\}_T$, where Θ_t are independent random vectors generated afresh for each tree, inducing randomness in both data sampling and feature selection [77].

Each tree h_t is trained on a bootstrap sample $D_t \subset D$ of size n , and at each internal node, a subset of $m \ll 10$ features is selected uniformly at random from the 10 available features to determine the optimal split. The split

criterion maximizes the information gain (or minimizes the Gini impurity) over the selected features.

Let $G(x, j, \tau)$ denote the Gini impurity reduction from splitting on feature j at threshold τ :

$$G(x, j, \tau) = H(D_{\text{parent}}) - \sum_k H(D_k), \quad (15)$$

where $H(D) = 1 - \sum_{c \in \{-1, +1\}} p_c^2$, and p_c is the proportion of class c in the node.

The tree recursively partitions the feature space X into axis-aligned regions $\{R_k\}_T$, assigning to each region a class label based on majority vote:

$$h_t(x) = \arg \max_{c \in \{-1, +1\}} \sum_{i: x_i \in R_k, y_i = c} w_i, \quad (16)$$

where w_i may be uniform or bootstrap-weighted.

The ensemble prediction is:

$$f_{\text{RF}}(x) = \text{sign} \left(\frac{1}{T} \sum_{t=1}^T h_t(x; \Theta_t) \right), \quad (17)$$

with probability estimates given by:

$$P(y = +1|x) = \frac{1}{T} \sum_{t=1}^T \mathbb{I}(h_t(x) = +1). \quad (18)$$

4) Extra Trees (Extremely Randomized Trees)

The Extra Trees ensemble enhances randomness by eliminating bootstrapping and introducing stochasticity in split thresholds [78]. Each tree h_{ext} is trained on the full dataset D , and at each node, a random subset of m' features is selected. For each candidate feature j , s random split points $\{\tau_{j,s'}\}_{s'}$ are drawn uniformly from the empirical range of feature j in the current node.

The split is chosen as:

$$(j^*, \tau^*) = \arg \max_{j \in F_{\text{rand}}, \tau \in \{\tau_{j,s'}\}} G(x, j, \tau), \quad (19)$$

where F_{rand} is the randomly selected feature subset.

5) AdaBoost (Adaptive Boosting)

AdaBoost constructs a strong classifier as a weighted linear combination of weak learners (typically decision stumps), iteratively reweighting the training instances to focus on previously misclassified samples [79].

Let $h_t : X \rightarrow \{-1, +1\}$ be the weak learner at iteration t , selected to minimize the weighted error:

$$\epsilon_t = \sum_{i=1}^n w_i^{(t)} \mathbb{I}(h_t(x_i) \neq y_i), \quad (20)$$

where $w_i^{(t)}$ is the weight of instance i at round t . The initial weights are uniform: $w_i^{(1)} = \frac{1}{n}$.

The coefficient α_t for the weak hypothesis is:

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right), \quad (22)$$

assuming $\epsilon_t < \frac{1}{2}$.

The final strong classifier is:

$$f_{\text{AdaBoost}}(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right). \quad (25)$$

6) CatBoost

CatBoost is a gradient boosting framework that handles categorical features natively and reduces prediction shift via ordered boosting. It constructs an additive model [80]:

$$f_M(x) = f_0 + \sum_{m=1}^M \gamma_m h_m(x), \quad (26)$$

where f_0 is the initial approximation (e.g., log-odds of class priors), h_m are weak learners (symmetric decision trees), and γ_m are coefficients determined by line search.

CatBoost employs an ordered boosting mechanism to mitigate overfitting. It creates a random permutation σ of the training indices and defines an “internal time” $t_i = \sigma(i)$ for each instance. The model at step m is trained on a simulated history, using only instances j with $t_j < t_i$ to compute gradients and construct trees.

The final prediction is:

$$f_{\text{CatBoost}}(x) = \text{sign}(f_M(x)), \quad (29)$$

with probabilities obtained via the logistic link:

$$P(y = +1|x) = \sigma(f_M(x)), \quad (30)$$

where $\sigma(z) = \frac{1}{1+e^{-z}}$ is the sigmoid function.

7) Feature Space and Model Integration

All models operate on the same preprocessed feature vector $x \in \mathbb{R}^{10}$, with categorical variables binarized and continuous ones standardized. The joint modeling strategy enables a comprehensive comparison of margin-based, ensemble, and boosting paradigms in diagnosing thyroid dysfunction from clinical markers. Hyperparameters for all models are optimized via nested cross-validation to avoid bias, ensuring statistical validity of performance estimates (See algorithm 1).

E. INTERPRETABILITY: SHAP AND COUNTERFACTUAL EXPLANATIONS

1) Motivation

Clinical decision support benefits from explanations at two levels: global—which predictors the model relies on across the cohort—and local—why a specific patient was classified as Normal, Hypothyroid, or Hyperthyroid. We therefore pair our classifiers with SHAP (additive feature attributions) and with counterfactual explanations that probe decision boundaries by asking what minimal, plausible change would alter a prediction.

F. COUNTERFACTUAL EXPLANATIONS

Given an input x with predicted class $\hat{y} = \arg \max_k f_k(x)$ and target class $y^* \neq \hat{y}$, we seek a counterfactual x^* that minimally perturbs x while flipping the decision:

$$x^* \in \arg \min_{x' \in F} (\|W(x' - x)\|_1 + \lambda \ell(f(x'), y^*)), \quad (35)$$

subject to the following constraints:

$$x'_j = x_j \quad \forall j \in \mathcal{X}, \quad (36)$$

$$l \leq x' \leq u, \quad (37)$$

where:

- $\|W(x' - x)\|_1$ is a weighted sparsity cost, with $W = \text{diag}(w_1, \dots, w_d)$ and $w_j = \frac{1}{\text{MAD}_j}$ to normalize units.
- $\ell(f(x'), y^*)$ is a loss function encouraging class y^* , e.g., $\ell = -\log P(y^*|x') + \max_{k \neq y^*} \log P(k|x')$.
- $\mathcal{X}$ indexes immutable features (e.g., sex, thyroid surgery); only labs {TSH, FTI, T3, TT4, T4U} vary.
- l, u are bounds from empirical percentiles (e.g., 1st–99th) to ensure plausibility.

We solve (49) via projected coordinate descent guided by local SHAP magnitudes, updating a mutable feature $j \notin \mathcal{X}$ by:

$$x' \leftarrow \Pi_{[l,u]} (x' - \eta \nabla_{x'} (\|W(x' - x)\|_1 + \lambda \ell(f(x'), y^*))), \quad (38)$$

until the prediction flips to y^* or a maximum number of iterations is reached.

We report:

- The scaled distance $\Delta = \|W(x^* - x)\|_1$ from (49);
- Per-feature deltas $\delta_j = x^* - x_j$ for all mutable $j \notin \mathcal{X}$;
- The decision shift $\Delta f_k = f_k(x^*) - f_k(x)$ for each class k .

These quantify how small, clinically bounded edits to thyroid labs can alter model diagnoses.

G. PERFORMANCE AND EVALUATION METRICS

1) Three-class Formulation

Let $D = \{(x_i, y_i)\}_N$ denote the test set, with $x_i \in \mathbb{R}^d$ and $y_i \in \{1, 2, 3\}$ corresponding to Normal, Hypothyroid, and Hyperthyroid, respectively. Define the one-hot label matrix $Y \in \{0, 1\}^{N \times 3}$, where $Y_{ik} = \mathbb{1}\{y_i = k\}$, and the model's class-probability matrix $P \in [0, 1]^{N \times 3}$, where $P_{ik} = P(y = k|x_i)$.

The predicted labels are $\hat{y}_i = \arg \max_k P_{ik}$, and the cross-entropy on the test set is:

$$L_{CE} = - \sum_{i=1}^N \sum_{k=1}^3 Y_{ik} \log P_{ik}. \quad (39)$$

2) Confusion Matrix and Derived Metrics

Let the confusion matrix $C \in \mathbb{N}^{3 \times 3}$ be defined with rows corresponding to true classes and columns corresponding to predicted classes:

$$C = \begin{pmatrix} C_{11} & C_{12} & C_{13} \\ C_{21} & C_{22} & C_{23} \\ C_{31} & C_{32} & C_{33} \end{pmatrix}$$

Overall accuracy is the normalized trace:

$$\text{Acc} = \frac{\text{tr}(C)}{N} = \frac{C_{11} + C_{22} + C_{33}}{N}. \quad (40)$$

For class $k \in \{1, 2, 3\}$, define the following metrics:

$$TP_k = C_{kk},$$

$$FN_k = \sum_{j \neq k} C_{kj},$$

$$FP_k = \sum_{i \neq k} C_{ik},$$

$$TN_k = \sum_{i,j \neq k} C_{ij}.$$

Then:

$$\text{Precision}_k = \frac{TP_k}{TP_k + FP_k}, \quad (41)$$

$$\text{Recall}_k = \frac{TP_k}{TP_k + FN_k}, \quad (42)$$

$$F1_k = 2 \cdot \frac{\text{Precision}_k \cdot \text{Recall}_k}{\text{Precision}_k + \text{Recall}_k}. \quad (43)$$

Let the per-class supports be $s_k = \sum_{i=1}^3 C_{ki}$ and $S = \sum_{k=1}^3 s_k = N$. We report macro and support-weighted averages:

$$\text{Precision}_{\text{macro}} = \frac{1}{3} \sum_{k=1}^3 \text{Precision}_k,$$

$$\text{Precision}_{\text{weighted}} = \frac{1}{N} \sum_{k=1}^3 s_k \cdot \text{Precision}_k,$$

$$\text{Recall}_{\text{macro}} = \frac{1}{3} \sum_{k=1}^3 \text{Recall}_k,$$

$$\text{Recall}_{\text{weighted}} = \frac{1}{N} \sum_{k=1}^3 s_k \cdot \text{Recall}_k,$$

$$F1_{\text{macro}} = \frac{1}{3} \sum_{k=1}^3 F1_k, \quad F1_{\text{weighted}} = \frac{1}{N} \sum_{k=1}^3 s_k \cdot F1_k. \quad (46)$$

H. OVR ROCAUC FROM THE PROBABILITY MATRIX

For one-vs-rest (OvR) ROC per class k , regard class k as positive and $\{1, 2, 3\} \setminus \{k\}$ as negative, using scores $\{P_{ik}\}_N$. The ROC curve is:

$$TPR(t) = \frac{1}{N} \sum_{i=1}^N \mathbb{I}\{y_i = k\} \mathbb{I}\{P_{ik} \geq t\},$$

$$FPR(t) = \frac{1}{N \sum_{i=1}^N \mathbb{I}\{y_i \neq k\}} \sum_{i=1}^N \mathbb{I}\{y_i \neq k\} \mathbb{I}\{P_{ik} \geq t\}.$$

The per-class AUC is:

$$\text{AUC}_k = \int_0^1 TPR_k^{-1}(u) \cdot FPR_k(u) du. \quad (48)$$

We summarize with:

$$\text{AUC}_{\text{macro}} = \frac{1}{3} \sum_{k=1}^3 \text{AUC}_k,$$

$$\text{AUC}_{\text{weighted}} = \frac{1}{N} \sum_{k=1}^3 s_k \cdot \text{AUC}_k.$$

III. RESULTS AND DISCUSSION

We evaluate thyroid status classification (Normal, Hypothyroid, Hyperthyroid) using the real UCI dataset (Section 3.2), our top-10 consensus features (TSH, FTI, T3, TT4, T4U, on thyroxine, query hypothyroid, psych, thyroid surgery, sex), and the nested cross-validation (CV) protocol in Algorithm 1. Unless otherwise stated, all summary metrics, confusion matrices, and ROC curves are computed from out-of-fold (OOF) predictions from the 5-fold outer CV, thereby reflecting performance strictly on unseen data within each fold.

Standardization was applied only to continuous laboratory variables (TSH, FTI, T3, TT4, T4U) within the inner training fold using StandardScaler. The fitted scaler was then used to transform validation and test splits, preventing data leakage. Binary/categorical features (e.g., sex, on thyroxine) were kept as $\{0,1\}$. Scaling was disabled for tree-based models (RF, CatBoost, etc.) since they do not require it; only SVC pipelines included scaling as a mandatory step.

Table 4 summarizes the nested cross-validation (CV) results for five classifiers under two feature regimes: the full feature set and the top-10 consensus features. Hyperparameters were optimized using grid search over predefined ranges with a 3-fold inner CV, targeting accuracy. Final models were retrained on the full outer training folds with the best parameters identified.

Across all models, accuracy exceeded 0.96 and AUC was 0.990, indicating a relatively separable three-class classification task. Ensemble methods performed particularly well: Random Forest and CatBoost consistently achieved the best mean scores across both feature regimes (Accuracy/F1 0.996–0.997; AUC 0.999–1.000), while SVC trailed slightly (Accuracy 0.965–0.967; AUC 0.990–0.992). Performance with the top-10 features was either preserved or marginally improved compared to the full feature set (e.g., Extra Trees: 0.981 $\rightarrow$ 0.989 accuracy; Random Forest: 0.996 $\rightarrow$ 0.997

TABLE 5: Nested CV on the Real Thyroid Dataset. Metrics are outer-fold means (3 d.p.).

Model	Feat.	Acc.	Rec.	Prec.	F1	AUC	Best params
SVC	Full	0.965	0.965	0.962	0.962	0.992	C=1, kernel=linear, gamma=scale
SVC	Top-10	0.967	0.967	0.965	0.963	0.990	
Random Forest	Full	0.996	0.996	0.996	0.996	0.999	n_estimators=100, max_depth=None, criterion=gini
Random Forest	Top-10	0.997	0.997	0.997	0.997	1.000	
Extra Trees	Full	0.981	0.981	0.981	0.980	0.998	n_estimators=200, max_depth=None, criterion=gini
Extra Trees	Top-10	0.989	0.989	0.989	0.989	0.999	
AdaBoost	Full	0.990	0.990	0.991	0.990	0.996	n_estimators=50, learning_rate=1.0
AdaBoost	Top-10	0.990	0.990	0.991	0.990	0.996	
CatBoost	Full	0.996	0.996	0.996	0.996	1.000	depth=4, iterations=200, learning_rate=0.1, l2_leaf_reg=3
CatBoost	Top-10	0.996	0.996	0.996	0.996	1.000	

accuracy), consistent with the hypothesis that a small set of physiologically salient labs and flags carry most of the predictive signal.

Searched hyperparameter ranges and best configurations:

- SVC was tuned over $C \in \{0.1, 1, 10, 100\}$, kernels {linear, rbf}, and $\gamma \in \{\text{scale}, \text{auto}\}$. The best-performing model used $C = 1$, kernel=linear, $\gamma = \text{scale}$.
- Random Forest was optimized over $n_{\text{estimators}} \in \{50, 100, 200\}$, $\text{max_depth} \in \{\text{None}, 10, 20\}$, and criterion {gini, entropy}. The best setting was $n_{\text{estimators}} = 100$, $\text{max_depth}=\text{None}$, criterion=gini.
- Extra Trees was searched over identical ranges to RF, with $n_{\text{estimators}} = 200$, $\text{max_depth}=\text{None}$, criterion=gini selected as optimal.
- AdaBoost was tuned with $n_{\text{estimators}} \in \{50, 100, 200\}$ and learning_rate {0.5, 1.0, 1.5}, converging on $n_{\text{estimators}} = 50$, learning_rate=1.0.
- CatBoost was optimized over depth {4, 6, 8}, iterations {100, 200}, learning_rate {0.03, 0.1, 0.3}, and $l2_{\text{leafreg}} \in \{1, 3, 5\}$. The best configuration was depth=4, iterations=200, learning_rate=0.1, $l2_{\text{leafreg}} = 3$.

Given this profile, Random Forest was designated as the primary model for reporting and deployment: it delivers state-of-the-art performance, produces well-calibrated probabilities, requires minimal preprocessing, and supports exact TreeSHAP explanations. CatBoost is retained as a near-identical, interpretable baseline.

The outstanding performance of Random Forest and CatBoost arises from their ability to capture nonlinear interactions among laboratory values while avoiding overfitting through ensemble averaging. Their hierarchical splitting structure aligns with clinical reasoning (e.g., "if TSH is elevated $\rightarrow$ check FTI"). By contrast, SVC, although competitive, is less effective in multi-class settings without kernel fine-tuning.

Finally, to examine the impact of class imbalance, all models were retrained using SMOTE-balanced data. As detailed in Appendix C (Table 19), minority-class recall improved modestly (<2% for hyperthyroid cases), while precision declined slightly due to increased false positives. Therefore, we retained the original imbalanced setup to reflect real-world prevalence and avoid synthetic artifacts.

IV. WHY THE TOP-10 FEATURE SET FOR THE UI AND DEPLOYMENT?

Although the full feature set yields excellent scores, the Top-10 consensus set (TSH, FTI, T3, TT4, T4U, on thyroxine, query hypothyroid, psych, thyroid surgery, sex) offers multiple advantages with no loss of accuracy:

- **Parsimony without performance loss:** As Table 4 shows, the Top-10 matches or slightly improves Accuracy/F1/AUC for the best models, suggesting that remaining variables are either redundant or weakly informative.
- **Clinically aligned signal:** The selected labs (TSH, FT4/FTI, TT4, T3, T4U) directly reflect hypothalamic pituitary-thyroid axis physiology; therapy flags (on thyroxine, thyroid surgery) encode iatrogenic states; sex is a known modifier. This improves face validity and trust.
- **Interpretability and XAI:** Fewer inputs yield crisper SHAP summaries and more actionable counterfactuals (Sec. 3.6), simplifying communication of "why" and "what-if" to clinicians.
- **Operational robustness:** The Top-10 fields are routinely available in EHRs, reducing missing-data risk, user burden, and latency at prediction time.
- **Regularization by design:** Restricting to the most stable features mitigates spurious correlations and potential shift from rare indicators; this complements statistical regularization used during training.

Figure 8 and 9 present, for each candidate model, the out-of-fold (OOF) confusion matrix (top row) alongside the corresponding multiclass ROC curves (bottom row). Across all tree and boosting methods—Random Forest, ExtraTrees, AdaBoost, and CatBoost—the ROC curves hug the upper-left corner, yielding micro and macro AUCs essentially at or extremely close to 1.000. The linear SVC performs strongly but consistently below the ensembles, with micro/macro AUCs around 0.994–0.996, indicating a mild loss of discrimination when class boundaries are non-linear. The confusion matrices corroborate these findings: most predictions lie on the diagonal, especially for the Normal class, while the few misclassifications that remain are predominantly Hyperthyroid (and to a lesser extent Hypothyroid) cases labeled as Normal, pointing to a tight decision boundary between these two phenotypes.




FIGURE 8: Out-of-fold (5-fold outer CV) confusion matrices and ROC curves for all models (SVC, Random Forest, Extra Trees, AdaBoost, CatBoost) using the real dataset and the chosen feature set. Rows show true classes and columns show predicted classes; values are aggregated over outer folds and Curves include per-class ROC along with micro- and macro-average ROC; AUCs are computed from OOF probabilities and aggregated across folds.

Among the ensembles, CatBoost and Random Forest exhibit the cleanest matrices, with vanishingly small off-diagonal counts and ROC curves that are effectively saturated. ExtraTrees achieves a similar AUC profile but shows slightly more spillover from Hyper/Hypo into Normal. AdaBoost also attains near-perfect AUCs, though its confusion matrix reveals a few additional Normal→Hyper,Hypo assignments, suggesting a somewhat sharper boundary that occasionally over-flags positives. SVC’s matrix shows the highest rate of Hyper→Normal errors among the models, which is consistent with the limitations of a linear separator in the presence of feature interactions.

Taken together, these results indicate that non-linear relations among the top clinical variables (notably TSH, FTI, T3, TT4, T4U, and key indicators such as on thyroxine and query hypothyroid) are critical for accurate triage across Normal, Hypothyroid, and Hyperthyroid. For deployment, CatBoost or Random Forest provide the best balance of sensitivity and specificity with minimal cross-class leakage, whereas SVC remains a useful linear baseline but risks missing Hyperthyroid presentations. In settings where the clinical cost of false negatives is asymmetric, threshold calibration on the probabilistic outputs of the tree/boosting models can further bias decisions toward safeguarding recall in the disordered classes without materially sacrificing overall performance.

Deep learning models were excluded due to the moderate size of the UCI thyroid dataset (7,200 samples), which increases the risk of overfitting without large-scale pretraining or extensive regularization. Moreover, our focus was on fully interpretable pipelines, where SHAP values can be computed exactly using TreeSHAP for tree ensembles.




FIGURE 9: Global feature importance based on mean absolute SHAP values. Features are ranked by their average contribution magnitude across all predictions. TSH overwhelmingly dominates the decision logic, followed by thyroxine, T3, TT4, and FTI. Remaining variables (T4U, thyroid surgery, sex, query hypothyroid, psych) contribute minimally.

Most deep networks require KernelSHAP or similar approximations, which are computationally expensive and less stable. Future work may explore lightweight neural architectures with integrated explainability or hybrid symbolic-AI frameworks. Including deep models would also complicate counterfactual generation due to non-convex loss landscapes and longer inference times—factors critical for clinical deployment feasibility.

A. INTERPRETABILITY

1) Explainable Artificial Intelligence (XAI): Mathematical Foundation and Clinical Validation

To ensure transparency and clinical fidelity, we ground our interpretability analysis in the axiomatic framework of SHAP (SHapley Additive exPlanations), which uniquely satisfies desirable properties such as local accuracy, missingness, and consistency [?]. Below, we present a formal derivation of the SHAP decomposition, validate its mathematical soundness, and provide clinical interpretations through structured analysis of representative patient cases.

This table bridges abstract mathematics and clinical reasoning. For example, if $\varphi_j(x) > 0$, it indicates that feature j pushes the model toward class k , while $\varphi_j(x) < 0$ suggests it detracts from it. The additive structure ensures no interaction terms are hidden, making the model fully decomposable.

The global SHAP analysis confirms that:

- **TSH** is the most influential feature by a wide margin, reflecting its central role in thyroid physiology and clinical diagnosis.
- **On thyroxine** and **T3/TT4/FTI** carry secondary importance, modulating model predictions depending on medication status and hormone levels.
- **T4U** provides minor contribution, aligning with its

TABLE 6: Mathematical structure and clinical interpretation of SHAP components

Equ	Mathematical Meaning	Clinical Interpretation
(1)	$f_k(x) = \varphi_k + \sum_{j=1}^d \varphi_j(x)$	Decomposes prediction into baseline and feature effects. Patient's predicted risk = population baseline + personalized adjustments from biomarkers.
(2)	$\varphi_0 = \mathbb{E}_{x \sim D}[f_k(x)]$	Expected model output for class k . Represents average risk of class k in the reference population.
(3)	$\varphi_j(x) = \sum_{S \subseteq N \setminus \{j\}} \frac{ S !(d- S -1)!}{d!} [f_k(x_S \cup \{j\}) - f_k(x_S)]$	Shapley value for marginal contribution of feature j . Quantifies how much each clinical variable contributes to diagnosis for class k .
(4)	$f_k(x_S) = \mathbb{E}_{\bar{x} \sim D}[f_k(x_S, \bar{x})]$	Conditional expectation over missing features. Model's estimated output when only some features are known (e.g., partial labs).
(5)	$\pi(S) = \frac{d!}{ S !(d- S)!}$	Kernel weight for subset S in KernelSHAP. Prioritizes balanced subsets to stabilize approximation.
(6)	$\bar{\Phi}(k) = \frac{1}{n_{\text{test}}} \sum_{i=1}^{n_{\text{test}}} \varphi_j(x_i) $	Average absolute SHAP value. Ranks, which features the most influential predictions for class k across all patients.

supportive but less diagnostic role in thyroid function assessment.

- **Sex, psych, and query hypothyroid** have near-zero importance, indicating the model avoids spurious or biased associations.

2) Proof of SHAP Decomposition via Shapley Values

We now sketch the derivation of Equation (30), showing that the Shapley value provides a unique, fair allocation of prediction differences. Let $v(S) = f_k(x_S)$ be the value function representing the model's expected output when only features in coalition $S \subseteq N = \{1, \dots, d\}$ are observed, and the rest are marginalized out conditionally:

$$v(S) = \mathbb{E}_{\bar{x}|x}[f_k(x)]$$

In cooperative game theory, each feature is a “player”, and $v(S)$ is the “payout” of coalition S . The Shapley value for player j is defined as the expected marginal contribution over all possible permutations of feature inclusion:

$$\varphi_j(x) = \frac{|S|!(d-|S|-1)!}{d!} \sum_{S \subseteq N \setminus \{j\}} [v(S \cup \{j\}) - v(S)]$$

It is a known result that this value is the unique solution satisfying the following axioms: 1. **Efficiency**: $\sum_{j=1}^d \varphi_j(x) = v(N) - v(\emptyset) = f_k(x) - \varphi_0$ 2. **Symmetry**: Interchangeable features get equal credit. 3. **Dummy**: If j adds no value, $\varphi_j = 0$. 4. **Additivity**: For combined games, values add. Therefore, the decomposition:

$$f_k(x) = v(\emptyset) + \sum_{j=1}^d \varphi_j(x)$$

is uniquely justified under these fairness criteria. This proves Equation (30) and validates the use of Shapley values for explainability.

TABLE 7: SHAP Decomposition for Patient 19 – Class 1 (Hypothyroid), $f_1(x_{19}) = 0.9900$

Feature j	Shapley Value $\varphi_1(x_{19})$
TSH	-0.0044
FTI	-0.0041
TT4	-0.0022
T3	-0.0005
T4U	+0.0002
on thyroxine	+0.0002
query hypothyroid	0
sex	0
psych	0
thyroid surgery	0
Sum $\Sigma \varphi_1(x_{19})$	Predicted $f_1(x_{19})$
	0.9900

TABLE 8: SHAP Decomposition for Patient 33 – Class 2 (Hyperthyroid), $f_0(x_{33}) = 0.9900$

Feature j	Shapley Value $\varphi_0(x_{33})$
TSH	-0.0500
on thyroxine	+0.0038
FTI	+0.0034
TT4	+0.0014
T4U	+0.0010
T3	-0.0003
query hypothyroid	+0.0001
sex	-0.0001
psych	0
thyroid surgery	0
Sum $\Sigma \varphi_0(x_{33})$	Predicted $f_0(x_{33})$
	0.9900

3) SHAP Decomposition for Representative Patients

4) Consistency Between Global Importance and Local Explanations

5) Interpretation of SHAP Values for Key Patients

Table 9 and Table 10 demonstrate consistency between global and local XAI: features with high global importance (e.g., TSH) dominate individual predictions, while irrelevant features contribute negligibly. Furthermore, the directional effects match clinical expectations:

- **TSH** $\uparrow \rightarrow$ favors Hypothyroid (Class 1)
- **TSH** $\downarrow \rightarrow$ favors Hyperthyroid (Class 2)

TABLE 9: SHAP Decomposition for Patient 0 – Class 3 (Normal), $f_2(x_0) = 1.0000$

Feature j	Shapley Value $\varphi_2(x_0)$
TSH	+0.0532
T4U	-0.0044
on thyroxine	-0.0038
T3	+0.0036
TT4	+0.0027
FTI	+0.0001
query hypothyroid	0
sex	0
psych	0
thyroid surgery	0
Sum $\Sigma \varphi_2(x_0)$	Predicted $f_2(x_0)$
	1.0000

TABLE 10: Consistency Between Global Importance and Local Explanations

Feature	Global $\Phi(k)$ Trend	Local Example
TSH	High across all classes	Index 19: -0.0044; Index 33: -0.0500; Index 0: +0.0532
FTI	Moderate global importance	Index 33: +0.0034 (Hyper)
on thyroxine	Low-to-moderate	Index 0: -0.0038 (Normal)
sex, psych	Near zero globally	All patients: 0

• **TSH mid-range** $\rightarrow$ favors Normal (Class 3)

Thus, the model learns interpretable, physiologically grounded rules, and SHAP faithfully reflects them.

We have formally established that: 1. The SHAP decomposition is mathematically rigorous, derived from the unique solution to a cooperative game. 2. Local explanations are clinically coherent, matching endocrinological knowledge. 3. Global and local views are consistent, ensuring scalability of insights.

These results validate our models not only in accuracy but also in epistemic trustworthiness, a prerequisite for deployment in clinical decision support systems.

6) Counterfactual Explainable AI

To move beyond static explanations like SHAP, we employ counterfactual explanations (CF) [83], which answer: "What minimal changes to the input would cause the model to change its prediction?" This is particularly valuable in clinical settings, where understanding how to alter a patient's status to reach a different diagnostic outcome can guide treatment planning.

Given an input x with predicted class $\hat{y} = \arg \max_k f_k(x)$, and a target class $y^* \neq \hat{y}$, we seek a counterfactual x^* that minimally perturbs x while flipping the decision. We solve:

$$x^* \in \arg \min_{x' \in F} (\|W(x' - x)\|_1 + \lambda \ell_1 f(x', y^*)) \quad (49)$$

subject to the constraints:

$$x'_j = x_j, \forall j \in \mathcal{X} \quad (50)$$

$$l \leq x' \leq u \quad (51)$$

Where:

- $\|W(x' - x)\|_1$ is a weighted L1 norm that promotes sparse, interpretable changes, with $W = \text{diag}(w_1, \dots, w_d)$ and $w_j = 1/\text{MAD}_j$ to normalize across units.
- $\ell_1 f(x', y^*) = -\log P(y^*|x') + \max_{k \neq y^*} \log P(k|x')$ encourages confidence in y^* while penalizing competing classes.
- $\mathcal{X}$ denotes the set of immutable features (e.g., sex, thyroid surgery) that remain fixed during optimization.
- l, u are lower and upper bounds (1st–99th percentiles) that ensure biologically plausible values.

We solve this via projected coordinate descent guided by local SHAP magnitudes, updating mutable features $j \notin \mathcal{X}$ as:

$$x' \leftarrow \Pi_{[l, u]} (x' - \eta \nabla_{x'} [\|W(x' - x)\|_1 + \lambda \ell_1 f(x', y^*)]) \quad (52)$$

until the prediction flips or convergence is reached.

The original patient has: TSH = 0.008 (very low; strong hyperthyroid signal), FTI = 0.100 (elevated), TT4 = 0.066 (normal/high-normal); other features are neutral/negative. The model predicts Class 2 (Hyperthyroid) with high confidence.

Now consider the three counterfactuals:

- **CF 1:** Decrease FTI from 0.100 to 0.03145 (68.5% reduction), keep others the same. Mimics low free hormone despite suppressed TSH—e.g., recovery or non-thyroidal illness.
- **CF 2:** Increase TT4 to 0.478 (650% increase) while slightly reducing FTI. Represents binding-protein abnormalities where total T4 is high but free fraction is low.
- **CF 3:** Keep labs similar but set query hypothyroid = 1. Suggests that a strong clinical suspicion flag could shift the classification.

All three counterfactuals flip the prediction to Class 1 (Hypothyroid), demonstrating:

- Sensitivity to FTI: the most consistent lever is sharply lowering FTI.
- Plausibility: changes respect biological bounds.
- Actionability: clinicians can read these as "what-if" scenarios (e.g., monitor FTI).

The scaled distance Δ is smallest for CF 1, indicating the most efficient path to flip the diagnosis—aligning with SHAP, where FTI is a top contributor in hyperthyroid patients.

This paper (last two rows) reports nested-CV means and includes AUC plus XAI (SHAP) and counterfactual analyses.

As summarized in Table 11, prior studies on thyroid disease prediction consistently report very high accuracies, often exceeding 95% across a variety of machine learning models such as kNN, Decision Trees, AdaBoost, and

TABLE 11: State-of-the-Art Models for Thyroid Disease Prediction. This paper (last two rows) reports nested-CV means and includes AUC plus XAI (SHAP) and counterfactual analyses.

Ref	Year	Dataset	Model	Results
[59]	2020	UCI new-thyroid dataset	Logistic Regression, Decision Tree, k-Nearest Neighbor	kNN: 96.875% accuracy; DT: 87.5% accuracy; LR: 81.25% accuracy
[60]	2022	UCI thyroid disease dataset	Random Forest (with MLFS feature selection)	RF: 0.99 accuracy, 0.99 F1 score
[61]	2024	Clinical thyroid disease dataset (1232 samples)	Stacking Ensemble (LR, SVM, KNN, DT, ANN)	99.9% accuracy, 99.9% sensitivity, 99.8% specificity
[62]	2023	UCI dataset thyroid	AdaBoost (with DE optimization)	Accuracy: 0.998
[63]	2023	UCI dataset thyroid	Random Forest (with LASSO)	Acc 99%, Prec 97%, Rec 92%, F1 95%
[64]	2024	UCI dataset thyroid	Ensemble ML (Hard Voting on RF and DT)	Acc 99.71%, Sens 100%
[65]	2022	UCI dataset thyroid	Decision Tree, Random Forest	DT: 99.5758% accuracy; RF: 99.3107% accuracy
[Ours-Full]	2025	UCI thyroid dataset (this study)	Best model: Random Forest (nested CV)	Acc 0.996, F1 0.996, AUC 0.999
[Ours-Top-10]	2025	UCI thyroid dataset (Top-10 features)	Best model: Random Forest (nested CV)	Acc 0.997, F1 0.997, AUC 1.000

Random Forest variants. However, many of these works evaluate performance primarily using overall accuracy, which is vulnerable to inflation under class imbalance and does not guarantee robust minority-class discrimination.

In contrast, our results (last two rows) build upon this foundation by adopting a nested cross-validation framework on a real clinical dataset. Using both the full feature set and a reduced top-10 feature set, our Random Forest models achieved 99.6–99.7% accuracy and F1-scores, with near-perfect AUC values (0.999–1.000). Beyond numerical performance, our study uniquely incorporates explainable AI through SHAP-based feature attribution and counterfactual analysis, enhancing interpretability and clinical relevance. Thus, while our predictive scores are competitive with or superior to prior benchmarks, the critical advancement lies in rigorous validation, balanced evaluation metrics, and explicit incorporation of interpretability techniques.

V. FUTURE WORKS AND LIMITATIONS

While this study demonstrates high performance and strong interpretability in thyroid disease classification, several limitations must be acknowledged to ensure responsible translation into clinical practice. First, our analysis relies on the UCI Thyroid Disease dataset—a widely used benchmark—but it has inherent constraints that affect generalizability. The cohort consists of 7,200 samples, with significant class imbalance: 6,666 "Normal", 368 "Hyperthyroid", and only 166 "Hypothyroid" cases. This skew risks model bias toward the majority class and may impair sensitivity for rarer conditions, despite our models achieving robust metrics without synthetic oversampling. Although we evaluated SMOTE as a mitigation strategy (Appendix C, Table 19), its use introduced trade-offs in precision and raised concerns about generating non-physiological lab profiles. Therefore, we retained the original distribution to reflect real-world prevalence, but future work should explore cost-sensitive

learning or threshold tuning in settings where missing hypothyroidism carries higher clinical risk. Second, the dataset lacks etiologic markers such as TRH, ACTH, or autoantibodies (e.g., TPOAb, TRAb), preventing differentiation between primary and secondary forms of hypo- or hyperthyroidism. As a result, our model performs high-level triage rather than subtype-specific diagnosis—a necessary simplification given data limitations, but one that restricts its utility in complex endocrine evaluations. Third, the absence of multimodal data—including ultrasound imaging, cytology, or histopathology—limits applicability to benign dysfunction and excludes critical areas like thyroid nodule assessment or cancer detection. Integrating imaging features via radiomics or vision-based AI could extend this framework to malignancy screening in future hybrid systems. Fourth, while SHAP and counterfactual explanations significantly enhance trust and clinical interpretability, they incur computational overhead that may hinder real-time deployment, especially in low-resource or point-of-care environments. KernelSHAP-based methods for non-tree models (e.g., SVC) are particularly expensive, and even TreeSHAP, though efficient, requires additional inference steps. Model distillation into lightweight surrogates or optimized explanation caching could mitigate latency in production systems.

Finally, potential demographic biases exist in the dataset: approximately 70% of patients are female ($p < 0.001$), reflecting both the higher population prevalence of thyroid disorders in women and possible sampling bias. While sex was not a dominant feature in SHAP analysis, the underrepresentation of male and diverse age/ethnic groups raises concerns about fairness and equity across populations. Without explicit debiasing strategies, model performance may degrade in underrepresented subgroups. Future Work To address these limitations, we propose several directions for future research:

- External validation on multi-center, longitudinal cohorts

to assess generalizability across healthcare systems and patient demographics.

- Integration of multimodal inputs (labs + imaging + EHR notes) using fusion architectures to enable more granular diagnostics, including subtype classification and malignancy prediction.
- Adoption of fairness-aware ML techniques—such as adversarial debiasing, reweighting, or stratified sampling—to ensure equitable performance across sex, age, and ethnicity.
- Development of efficient XAI pipelines through model compression, surrogate modeling, or hardware-aware optimization for edge-device deployment.
- Exploration of federated learning frameworks to train on distributed electronic health records (EHRs) while preserving patient privacy and regulatory compliance.
- Design of clinician-in-the-loop interfaces that present SHAP values and counterfactuals in intuitive formats (e.g., dashboards, alerts), supporting shared decision-making.

By addressing these challenges, future iterations of this pipeline can move beyond retrospective benchmarking toward prospective, trustworthy, and deployable clinical decision support tools for endocrinology

VI. CONCLUSION

This study presents a robust, interpretable machine learning framework for thyroid disease classification that achieves state-of-the-art performance while providing clinically meaningful explanations. By applying rigorous nested cross-validation on the UCI Thyroid dataset, we evaluated five diverse classifiers—SVC, Random Forest, Extra Trees, AdaBoost, and CatBoost—and found that ensemble methods, particularly Random Forest and CatBoost, achieve near-perfect accuracy (up to 0.997), F1-score (0.997), and AUC (1.000) under both full and reduced feature regimes. A key contribution is the derivation of a TOP-10 consensus feature set—TSH, FTI, T3, TT4, T4U, on thyroxine, query hypothyroid, psych, thyroid surgery, and sex—selected through multi-criteria agreement across six feature ranking methods. Crucially, using only these ten physiologically grounded features not only simplifies the model but also slightly improves performance over the full set, demonstrating that high diagnostic fidelity can be maintained with minimal, actionable inputs. While the choice of "Top-10" was based on convergence of consensus rankings, future work should explore performance stability across varying subset sizes (e.g., Top-5 vs. Top-15) to further optimize parsimony. To bridge the gap between high performance and clinical trust, we integrated dual explainable AI (XAI) strategies: SHAP for global and local feature attribution, and counterfactual analysis to generate actionable “what-if” scenarios. These explanations align closely with endocrine physiology—for example, elevated TSH driving hypothyroid predictions or suppressed TSH with high FTI indicating hyperthyroidism—enabling clinicians to validate model logic and support patient

communication. Despite these strengths, several limitations must be acknowledged. The UCI dataset has class imbalance, lacks etiological markers (e.g., TPOAb, TRAb) for subtyping primary vs. secondary disease, and excludes multimodal data such as ultrasound. Additionally, SHAP and counterfactuals introduce computational overhead that may hinder real-time deployment in resource-constrained settings, and demographic imbalances (e.g., higher female representation) raise fairness concerns.

Future work will focus on external validation across multi-site cohorts, integration of multimodal data (labs + imaging + EHR), optimization of XAI for edge-device deployment, and implementation of fairness-aware algorithms to ensure equitable performance across populations. Federated learning frameworks could enable privacy-preserving training on distributed health records. Our pipeline combines methodological rigor—with leakage-free evaluation, consensus-based feature selection, and principled hyperparameter tuning—with dual-layer interpretability to create a trustworthy foundation for AI-assisted thyroid triage. This approach not only enhances diagnostic accuracy but also supports transparency, accountability, and clinician adoption—critical steps toward real-world deployment in endocrinology.

APPENDIX

A.1 ADDITIONAL FIGURES

This section presents supplementary exploratory data analysis (EDA) figures and visualizations that support the findings described in the main text. These include extended class distribution plots, boxplots for continuous variables, and additional bar plots for categorical variables. Figures are provided here to enhance transparency and reproducibility.




FIGURE 10: Class distribution of thyroid disease categories, including hyperthyroid, hypothyroid, and their respective subtypes.




FIGURE 11: Boxplots of selected continuous clinical features, illustrating variability and outliers across thyroid disease classes.




FIGURE 13: Correlation heatmap of clinical and demographic variables in the dataset.



FIGURE 12: Bar plots of categorical variables showing distributions across thyroid disease classes.




FIGURE 14: Patient 19: Waterfall plot and force plots. (a) SHAP decomposition for Patient 19 — predicted Class - 1 (Hypothyroid).

A.2 BASELINE CHARACTERISTICS

Table 12 shows the baseline characteristics by thyroid class in the SMOTE-balanced dataset.

B. FEATURE SELECTION

C. SMOTE-BALANCED THYROID DATASET MODELS PERFORMANCE

D. PATIENT-WISE INTERPRETATIONS

This appendix provides local (per-patient) explanations of model predictions using SHAP. For each patient, we show a waterfall plot (feature-wise contribution to the logit for the predicted class) and a force plot with panels for all three classes. The predicted class and its probability are printed on each force-plot panel.

REFERENCES

- [1] L.E. Braverman and D.S. Cooper, *Werner and Ingbar's The Thyroid: A Fundamental and Clinical Text*, 10th ed. Lippincott Williams & Wilkins, 2012.
- [2] R.H. Williams, *Williams Textbook of Endocrinology*, 13th ed. Elsevier, 2016.
- [3] M.P.J. Vanderpump, "The epidemiology of thyroid disease," *Br Med Bull*, vol. 99, no. 1, pp. 39–51, 2011.
- [4] J.R. Garber, R.H. Cobin, H. Gharib, et al., "Clinical practice guidelines for hypothyroidism in adults," *Thyroid*, vol. 22, no. 12, pp. 1200–1235, 2012.
- [5] B.R. Haugen, E.K. Alexander, K.C. Bible, et al., "2015 American Thyroid Association Management Guidelines for Adult Patients," *Thyroid*, vol. 26, no. 1, pp. 1–133, 2016.
- [6] D.S. Ross, H.B. Burch, D.S. Cooper, et al., "2016 American Thyroid Association Guidelines for Hyperthyroidism," *Thyroid*, vol. 26, no. 10, pp. 1343–1421, 2016.
- [7] M. Islam, R. Ferdousi, S. Rahman, "Machine learning approaches for thyroid disease diagnosis," *Comput Biol Med*, vol. 125, p. 103990, 2020.

TABLE 12: Baseline Characteristics by Thyroid Class (SMOTE-balanced)

Variable	Hypothyroid (N=6666)	Hyperthyroid (N=6666)	Normal (N=6666)	Overall (N=19998)	P-value
Age (years) < 0.001	Mean (SD)	0.540 (0.179)	0.557 (0.181)	0.518 (0.189)	0.538 (0.184)
	Median [Min, Max]	0.568 [0.0100, 0.890]	0.595 [0.0100, 0.910]	0.540 [0.0100, 0.970]	0.570 [0.0100, 0.970]
Sex < 0.001	No	5379 (80.7%)	5250 (78.8%)	4590 (68.9%)	15219 (76.1%)
	Yes	1287 (19.3%)	1416 (21.2%)	2076 (31.1%)	4779 (23.9%)
On thyroxine < 0.001	No	6040 (90.6%)	6666 (100%)	5743 (86.2%)	18449 (92.3%)
	Yes	626 (9.4%)	0 (0%)	923 (13.8%)	1549 (7.7%)
Query thyroxine < 0.001	No	6666 (100%)	6600 (99.0%)	6560 (98.4%)	19826 (99.1%)
	Yes	0 (0%)	66 (1.0%)	106 (1.6%)	172 (0.9%)
On antithyroid medication < 0.001	No	6665 (100%)	6656 (99.9%)	6578 (98.7%)	19899 (99.5%)
	Yes	1 (0.0%)	10 (0.2%)	88 (1.3%)	99 (0.5%)
Sick < 0.001	No	6665 (100%)	6474 (97.1%)	6407 (96.1%)	19546 (97.7%)
	Yes	1 (0.0%)	192 (2.9%)	259 (3.9%)	452 (2.3%)
Pregnant < 0.001	No	6666 (100%)	6666 (100%)	6588 (98.8%)	19920 (99.6%)
	Yes	0 (0%)	0 (0%)	78 (1.2%)	78 (0.4%)
Thyroid surgery < 0.001	No	6595 (98.9%)	6666 (100%)	6569 (98.5%)	19830 (99.2%)
	Yes	71 (1.1%)	0 (0%)	97 (1.5%)	168 (0.8%)
I-131 treatment < 0.001	No	6610 (99.2%)	6567 (98.5%)	6557 (98.4%)	19734 (98.7%)
	Yes	56 (0.8%)	99 (1.5%)	109 (1.6%)	264 (1.3%)
Query hypothyroid < 0.001	No	5761 (86.4%)	5795 (86.9%)	6269 (94.0%)	17825 (89.1%)
	Yes	905 (13.6%)	871 (13.1%)	397 (6.0%)	2173 (10.9%)
Query hyperthyroid < 0.001	No	6361 (95.4%)	6276 (94.1%)	6203 (93.1%)	18840 (94.2%)
	Yes	305 (4.6%)	390 (5.9%)	463 (6.9%)	1158 (5.8%)
Lithium < 0.001	No	6666 (100%)	6636 (99.6%)	6581 (98.7%)	19883 (99.4%)
	Yes	0 (0%)	30 (0.5%)	85 (1.3%)	115 (0.6%)
Goitre < 0.001	No	6666 (100%)	6666 (100%)	6607 (99.1%)	19939 (99.7%)
	Yes	0 (0%)	0 (0%)	59 (0.9%)	59 (0.3%)
Tumor < 0.001	No	6605 (99.1%)	6535 (98.0%)	6494 (97.4%)	19634 (98.2%)
	Yes	61 (0.9%)	131 (2.0%)	172 (2.6%)	364 (1.8%)
Hypopituitary = 1	No	6666 (100%)	6666 (100%)	6665 (100.0%)	19997 (100.0%)
	Yes	0 (0%)	0 (0%)	1 (0.0%)	1 (0.0%)
Psych < 0.001	No	6665 (100.0%)	6577 (98.7%)	6325 (94.9%)	19567 (97.8%)
	Yes	1 (0.0%)	89 (1.3%)	341 (5.1%)	431 (2.2%)

- [8] S. Duggal, R. Shukla, "Comparative analysis of machine learning classifiers for thyroid disease detection," *Procedia Comput Sci*, vol. 152, pp. 60–67, 2019.
- [9] T. Alyas, M. Iqbal, F. Khan, et al., "Classification of thyroid disease using ultrasound and machine learning," *Appl Sci*, vol. 11, no. 12, p. 5409, 2021.
- [10] S. Sonuç, T. Al-Dabbagh, S. Ismael, "Classification of thyroid diseases with machine learning models in Iraqi cohort," *Inform Med Unlocked*, vol. 29, p. 100913, 2022.
- [11] C. Molnar, *Interpretable Machine Learning*, 2nd ed. Lulu.com, 2020.
- [12] R. Guidotti, A. Monreale, S. Ruggieri, et al., "A survey of methods for explaining black box models," *ACM Comput Surv*, vol. 51, no. 5, pp. 1–42, 2018.
- [13] Q.A. Naseer and P. Jianbo, "Recent Advances in the Diagnosis, Pathophysiology, and Management of Thyroid Disorders," 2025.
- [14] S.L. Chew and R.D.G. Leslie, *Clinical endocrinology and diabetes: an illustrated colour text*, Elsevier Health Sciences, 2005.
- [15] C.A. Spencer, "Assay of thyroid hormones and related substances," 2015.
- [16] N. Mahmud, D.S. Goldberg, T. Bittermann, "Best practices in large database clinical epidemiology research in hepatology: barriers and opportunities," *Liver Transplantation*, vol. 28, no. 1, pp. 113–122, 2022.
- [17] B.J. Maron, M.Y. Desai, R.A. Nishimura, et al., "Diagnosis and evaluation of hypertrophic cardiomyopathy: JACC state-of-the-art review," *Journal of the American College of Cardiology*, vol. 79, no. 4, pp. 372–389, 2022.
- [18] M.J. Kofler, L.N. Irwin, E.F. Soto, et al., "Executive functioning heterogeneity in pediatric ADHD," *Journal of Abnormal Child Psychology*, vol. 47, no. 2, pp. 273–286, 2019.
- [19] M.G. Hanna, L. Pantanowitz, R. Dash, et al., "Future of artificial intelligence (AI)-machine learning (ML) trends in pathology and medicine," *Modern Pathology*, p. 100705, 2025.
- [20] M. Aftab, F. Mehmood, C. Zhang, et al., "AI in oncology: transforming cancer detection through machine learning and deep learning applications," *arXiv preprint arXiv:2501.15489*, 2025.
- [21] S. Rani, R. Kumar, B.S. Panda, et al., "Machine Learning-Powered Smart Healthcare Systems in the Era of Big Data: Applications, Diagnostic Insights, Challenges, and Ethical Implications," *Diagnostics*, vol. 15, no. 15, p. 1914, 2025.
- [22] BioRender, "Thyroid Gland Anatomy and Histology," [Online image]. Available: <https://www.biorender.com/template/thyroid-gland-anatomy-and-histology>, accessed Sept. 2, 2025.
- [23] S.D. Leo, S.Y. Lee, L.E. Braverman, "Hyperthyroidism," *The Lancet*, vol.

TABLE 13: Top 15 Features by Mutual Information

Feature	MI Score
TSH	0.770697
FTI	0.646006
T3	0.459210
TT4	0.401993
T4U	0.283670
on thyroxine	0.177822
query hypothyroid	0.139087
thyroid surgery	0.122216
psych	0.096831
sex	0.089303
tumor	0.084968
on antithyroid medication	0.066226
goitre	0.060244
query on thyroxine	0.048019
hypopituitary	0.027501

TABLE 14: Top 15 Features by ANOVA F-test

Feature	F Score	p-value
TSH	16130.300	0.000000e+00
FTI	3736.425	0.000000e+00
T3	3275.176	0.000000e+00
TT4	1509.104	0.000000e+00
T4U	105.026	1.344630e-23
on thyroxine	77.233	1.352180e-18
query hypothyroid	55.301	1.166261e-13
psych	47.133	1.000295e-11
thyroid surgery	37.066	1.128792e-09
sex	35.988	2.151339e-09
tumor	27.456	1.563441e-07
goitre	21.841	2.962861e-06
on antithyroid medication	18.327	1.006048e-05
query on thyroxine	12.627	3.604394e-04
hypopituitary	9.846	1.759224e-03

TABLE 15: Top 15 Features by Chi-Squared Test

Feature	Chi2 Score	p-value
TSH	12797.198	0.000000e+00
FTI	4840.842	0.000000e+00
T3	3921.280	0.000000e+00
TT4	2092.279	0.000000e+00
T4U	284.081	0.000000e+00
on thyroxine	165.426	2.603314e-37
query hypothyroid	124.101	4.183058e-28
thyroid surgery	100.792	2.893400e-23
psych	81.313	2.694059e-19
sex	66.139	4.401556e-16
tumor	53.774	2.174674e-13
on antithyroid medication	43.882	3.325129e-11
goitre	36.910	1.139160e-09
query on thyroxine	28.531	6.239801e-08
hypopituitary	17.646	2.653220e-05

TABLE 16: Top 15 Features by L1-Regularized Logistic Regression

Feature	L1 Coef
TSH	1.003326
FTI	0.802170
T3	0.634576
TT4	0.518217
T4U	0.372863
on thyroxine	0.290519
query hypothyroid	0.243778
psych	0.201179
thyroid surgery	0.185770
sex	0.174319
tumor	0.146719
on antithyroid medication	0.122891
goitre	0.108202
query on thyroxine	0.091370
hypopituitary	0.073456

388, no. 10047, pp. 906–918, 2016.

- [24] D.S. Ross, et al., "2016 American thyroid association guidelines for diagnosis and management of hyperthyroidism and other causes of thyrotoxicosis," *Thyroid*, vol. 26, no. 10, pp. 1343–1421, 2016.
- [25] D.Y. Gaitonde, K.D. Rowley, L.B. Sweeney, "Hypothyroidism: an update," *Am. Fam. Phys.*, pp. 244–251, 2012.
- [26] M.P. Vanderpump, "Epidemiology of thyroid disease," in *Encyclopedia of Endocrine Diseases*, pp. 486–495, 2018.
- [27] P.N. Taylor, et al., "Global epidemiology of hyperthyroidism and hypothyroidism," in *Yearbook of Paediatric Endocrinology*, 2018.
- [28] G. Medeiros-Neto, R.Y. Camargo, E.K. Tomimori, "Approach to and treatment of goiters," *Med. Clin. N. Am.*, vol. 96, no. 2, pp. 351–368, 2012.
- [29] D. Dauksiene, et al., "Factors associated with the prevalence of thyroid nodules and goiter in middle-aged euthyroid subjects," *Int. J. Endocrinol.*, 2017, 1–8.
- [30] M. Mesele, G. Degu, H. Gebrehiwot, "Prevalence and associated factors of goiter among rural children aged 6–12 years old in Northwest Ethiopia, cross-sectional study," *BMC Pub. Health*, vol. 14, no. 1, p. 130, 2014.
- [31] L.B. Sweeney, C. Stewart, D.Y. Gaitonde, "Thyroiditis: an integrated approach," *Am. Fam. Phys.*, vol. 90, no. 6, pp. 389–396, 2014.
- [32] A.P. Weetman and A.M. McGregor, "Autoimmune thyroid disease: developments in our understanding," *Endocrine Reviews*, vol. 5, no. 2, pp. 309–355, 1984.
- [33] S. Parangi and R. Phitayakorn, *Thyroid Disease*, Bloomsbury Publishing USA, 2010.
- [34] M. Sakr, *Midline Neck Swellings: Diagnostic and Therapeutic Challenges*, Springer Nature, 2024.
- [35] V. Nosé, *Diagnostic Pathology: Endocrine E-Book*, Elsevier Health Sciences, 2019.
- [36] S. Anerudh, "To Compare Accuracy of FNAC over Open Biopsy in Assessment of Cause of Nonthyroidal Neck Swelling," Master's thesis, Rajiv Gandhi University of Health Sciences, India, 2018.
- [37] A. Raviteja, "Clinicopathological Study of Carcinoma Thyroid With

Special Reference to Ultrasound Findings," Master's thesis, Rajiv Gandhi University of Health Sciences, India, 2015.

- [38] Y.E. Nikiforov, "Thyroid tumors: classification, staging and general considerations," in *Diagnostic Pathology and Molecular Genetics of the Thyroid*, Lippincott Williams Wilkins, Baltimore, 2012, pp. 108–118.
- [39] M.-R. Kwon, J. Shin, H. Park, et al., "Radiomics study of thyroid ultrasound for predicting BRAF mutation in papillary thyroid carcinoma: preliminary results," *Am. J. Neuroradiol.*, vol. 41, no. 4, pp. 700–705, 2020.
- [40] D. Chen, J. Hu, M. Zhu, et al., "Diagnosis of thyroid nodules for ultrasonographic characteristics indicative of malignancy using random forest," *BioData Mining*, vol. 13, no. 1, pp. 1–21, 2020.
- [41] A.J. Idaraga, G. Luong, V. Hsiao, et al., "False negative rates in benign thyroid nodule diagnosis: Machine learning for detecting malignancy," *J. Surg. Res.*, vol. 268, pp. 562–569, 2021.
- [42] S. Azia, P. Siva Kumar, A.S. Rao, "Machine learning techniques for thyroid disease diagnosis: a systematic review," in *Modern Approaches in Machine Learning and Cognitive Science: A Walkthrough*, pp. 203–212, 2020.
- [43] L. Aversano, M.L. Bernardi, M. Cimitile, et al., "Thyroid disease treatment prediction with machine learning approaches," *Procedia Comput. Sci.*, vol. 192, pp. 1031–1040, 2021.
- [44] M. Lomana, A.G. Weber, B. Birk, et al., "In silico models to predict the perturbation of molecular initiating events related to thyroid hormone homeostasis," *Chem. Res. Toxicol.*, vol. 34, no. 2, pp. 396–411, 2020.
- [45] R. Chaganti, F. Rustam, I. De La Torre Díez, et al., "Thyroid disease prediction using selective features and machine learning techniques," *Cancers*, vol. 14, no. 16, p. 3914, 2022.
- [46] D.C. Yadav and S. Pal, "Thyroid prediction using ensemble data mining techniques," *Int. J. Inf. Technol.*, vol. 14, no. 3, pp. 1273–1283, 2022.
- [47] L. Aversano et al., "Thyroid disease treatment prediction with machine learning approaches," *Procedia Comput. Sci.*, vol. 192, p. 1031–1040, 2021.

TABLE 17: Top 15 Features by Recursive Feature Elimination (RFE)

Feature	RFE Ranking
TSH	1
FTI	1
T3	1
TT4	1
T4U	1
on thyroxine	2
query hypothyroid	3
psych	4
thyroid surgery	5
sex	6
tumor	7
on antithyroid medication	8
goitre	9
query on thyroxine	10
hypopituitary	11

TABLE 18: Top 15 Features by Random Forest Importance

Feature	Importance
TSH	0.261223
FTI	0.197483
T3	0.152348
TT4	0.128215
T4U	0.091203
on thyroxine	0.041915
query hypothyroid	0.034782
psych	0.028611
thyroid surgery	0.025003
sex	0.021947
tumor	0.017430
on antithyroid medication	0.014622
goitre	0.012880
query on thyroxine	0.010172
hypopituitary	0.008980

TABLE 19: Nested CV on the SMOTE-balanced thyroid dataset. Metrics are outer-fold means (3 d.p.).

Model	Acc.	Rec.	Prec.	F1	AUC	Best params
SVC	0.987	0.987	0.987	0.987	0.998	C=1, kernel=linear
Random Forest	0.999	0.999	0.999	0.999	1.000	n=100, depth=None
Extra Trees	0.997	0.997	0.997	0.997	1.000	n=200, depth=None
AdaBoost	0.991	0.991	0.992	0.991	0.996	n=50
CatBoost	0.998	0.998	0.998	0.998	1.000	depth=4, iters=100




FIGURE 15: (b) Force plot for Patient 33 across all classes (panels show class-specific probabilities).

[48] X. Li, R. Shan, J. Chen, et al., "Development and validation of a three-stage, machine learning-based model to optimize the pre-operational clinical decision-making for the screening, diagnosis, and treatment of follicular thyroid neoplasms," 2024.

[49] H. Abbad Ur Rehman, C.-Y. Lin, Z. Mushtaq, et al., "Performance analysis of machine learning algorithms for thyroid disease," *Arab J. Sci. Eng.*, vol. 46, no. 10, pp. 9437–9449, 2021.

[50] I.M.D. Maysanjaya, H.A. Nugroho, N.A. Setiawan, "A comparison of classification methods on diagnosis of thyroid diseases," in *2015 International Seminar on Intelligent Technology and Its Applications (ISITIA)*, 2015.

[51] W. Ahmad, A. Ahmad, C. Lu, et al., "A novel hybrid decision support system for thyroid disease forecasting," *Soft Comput.*, vol. 22, no. 16, pp. 5377–5383, 2018.

[52] Chaganti, R., Rustam, F., De La Torre Díez, I., Mazón, J. L. V., Rodríguez, C. L., & Ashraf, I. (2022). Thyroid Disease Prediction Using Selective Features and Machine Learning Techniques. *Cancers*, 14(16), 3914. <https://doi.org/10.3390/cancers14163914>

[53] K. Salman, E. Sonuç, "Thyroid disease classification using machine learning algorithms," *J. Phys. Conf. Ser.*, vol. 1963, no. 1, p. 012140, 2021.

[54] M. Hu, et al., "Development and preliminary validation of a machine learning system for thyroid dysfunction diagnosis based on routine laboratory tests," *Commun. Med.*, vol. 2, no. 1, p. 9, 2022.

[55] S. Sankar, A. Potti, G.N. Chandrika, et al., "Thyroid disease prediction using XGBoost algorithms," *J. Mob. Multimed.*, vol. 18, no. 3, pp. 1–8, 2022.

[56] S.N. Kulkarni, A.R. Karwankar, "Thyroid disease detection using modified fuzzy hyperline segment clustering neural network," *Int. J. Comput. Technol.*, vol. 3, no. 3b, pp. 466–469, 2012.

[57] A. Keleş, A. Keleş, "ESTDD: expert system for thyroid diseases diagnosis," *Expert Syst. Appl.*, vol. 34, no. 1, pp. 242–246, 2008.

[58] A. Gjecka and M. Fetaji, "Particle Swarm Optimization-Based Artificial Neural Network for prediction of thyroid disease," 2024 13th

Mediterranean Conference on Embedded Computing (MECO), Budva, Montenegro, 2024, pp. 1–4, doi: 10.1109/MECO62516.2024.10577948.

[59] G. Chaubey, D. Bisen, S. Arjaria, et al., "Thyroid disease prediction using machine learning approaches," *National Academy Science Letters*, vol. 44, no. 3, pp. 233–238, 2021.

[60] R. Chaganti, F. Rustam, I. De La Torre Díez, et al., "Thyroid disease prediction using selective features and machine learning techniques," *Cancers*, vol. 14, no. 16, p. 3914, 2022.

[61] G. Obaido, O. Achilonu, B. Ogbuokiri, et al., "An improved framework for detecting thyroid disease using filter-based feature selection and stacking ensemble," *IEEE Access*, vol. 12, pp. 89098–89112, 2024.

[62] P. Gupta, F. Rustam, K. Kanwal, et al., "Detecting thyroid disease using optimized machine learning model based on differential evolution," *International Journal of Computational Intelligence Systems*, vol. 17, no. 1, p. 3, 2024.

[63] A. Sultana and R. Islam, "Machine learning framework with feature selection approaches for thyroid disease classification and associated risk factors identification," *Journal of Electrical Systems and Information Technology*, vol. 10, no. 1, p. 32, 2023.

[64] K.M.M. Uddin, A. Al Mamun, A. Chakrabarti, et al., "An ensemble machine learning-based approach to predict thyroid disease using hybrid feature selection," *Biomedical Analysis*, vol. 1, no. 3, pp. 229–239, 2024.

[65] K. Guleria, S. Sharma, S. Kumar, et al., "Early prediction of hypothyroidism and multiclass classification using predictive machine learning and deep learning," *Measurement: Sensors*, vol. 24, p. 100482, 2022.

[66] E. Ostrom, *Understanding Institutional Diversity*, Princeton University Press, 2009.

[67] N. Chen, J. Zhu, and E. Xing, "Predictive subspace learning for multi-view data: a large margin approach," in *Advances in Neural Information Processing Systems*, vol. 23, 2010.

[68] R. Bey, R. Goussault, F. Grolleau, et al., "Fold-stratified cross-validation for unbiased and privacy-preserving federated learning," *Journal of the American Medical Informatics Association*, vol. 27, no. 8, pp. 1244–1251, 2020.

[69] A. Seraj, M. Mohammadi-Khanaposthane, R. Daneshfar, et al., "Cross-validation," in *Handbook of Hydroinformatics*, pp. 89–105, Elsevier, 2023.

[70] D.K. Rakesh and P.K. Jana, "A general framework for class label specific mutual information feature selection method," *IEEE Transactions on Information Theory*, vol. 68, no. 12, pp. 7996–8014, 2022.




FIGURE 16: Patient 19: Waterfall and force plots. (a) SHAP decomposition for Patient 33 — predicted Class - 2 (Hypothyroid).





FIGURE 17: Patient 33: Waterfall and force plots. (b) Force plot for Patient 33 across all classes (panels show class-specific probabilities).

- [71] L.C. Parra, S. Haufe, and J.P. Dmochowski, "Correlated components analysis—extracting reliable dimensions in multivariate data," *arXiv preprint arXiv:1801.08881*, 2018.
- [72] A.G. Desai, "A STATISTICAL STUDY OF ANALYSIS OF INDIA GROSS DOMESTIC PRODUCT," *Statistics and Data Science*, vol. 145, 2023.
- [73] M.I. Muhammad, "MACHINE (SVM) RECURSIVE FEATURE REMOVAL," *Generations*, vol. 7, p. 9, 2023.
- [74] X. Li, Y. Wang, and R. Ruiz, "A survey on sparse learning models for feature selection," *IEEE Transactions on Cybernetics*, vol. 52, no. 3, pp. 1642-1660, 2020.
- [75] I. Altaf and M.A. Chachoo, "Permuted Gini Importance—PaP Impurity Measurement for Tree-Based Models," *Journal of Soft Computing and Data Mining*, vol. 5, no. 2, pp. 96-109, 2024.
- [76] O. Chapelle, "Support Vector Machines: principes d'induction, réglage automatique et connaissances a priori," *Knowledge Creation Diffusion Utilization*, 2004.
- [77] J. Ali, R. Khan, N. Ahmad, and I. Maqsood, "Random forests and decision trees," *International Journal of Computer Science Issues (IJCSI)*, vol. 9, no. 5, p. 272, 2012.
- [78] G. Kunapuli, *Ensemble Methods for Machine Learning*, Simon and Schuster, 2023.
- [79] A.J. Ferreira and M.A. Figueiredo, "Boosting algorithms: A review of methods, theory, and applications," in *Ensemble Machine Learning*, pp. 35-85, 2012.
- [80] *CatBoost is a gradient boosting framework that handles categorical features natively and reduces prediction shift via ordered boosting*, 2021.

- [81] J. Yang, "Fast TreeSHAP: Accelerating SHAP value computation for trees," *arXiv preprint arXiv:2109.09847*, 2021.
- [82] S.L. Chau, R. Hu, J. Gonzalez, and D. Sejdinovic, "RKHS-SHAP: Shapley values for kernel methods," in *Advances in Neural Information Processing Systems*, vol. 35, pp. 13050-13063, 2022.
- [83] E. Albini, J. Long, D. Dervovic, and D. Magazzeni, "Counterfactual Shapley additive explanations," in *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, pp. 1054-1070, 2022.



MOHAMMAD ZAHIDUL ALAM is a graduate student in Westcliff University's Master of Science in Information Technology program. As an Information Technology researcher, he focuses on cutting-edge advancements in artificial intelligence (AI), machine learning (ML), robotics technology, and business technology. Mohammad Zahidul Alam aspires to explore new areas in information technology. Since his academic career, he has been dedicated to achieving excellence and actively contributing to developing technologies that address real-world challenges. His research aims to develop innovative solutions that bridge technology and business applications, driving efficiency and growth in various industries.



RUKSHANDA RAHMAN is an IT Specialist with extensive experience in telecommunications, software engineering, and academic instruction. She earned her Master of Science in Computer Science from Westcliff University, Irvine, California, in 2024. Her professional background spans roles as Senior Software Engineer at Link3 Technologies Ltd., Junior Application Developer at EGeneration, and IT Specialist at MBI Health Services. She has also contributed to academia as a Teaching Assistant at Westcliff University and BRAC University. Recognized for her impactful contributions to technology, she received the FUELD Conference Award in Las Vegas in 2025. Her research and professional interests include system design, software development, and emerging technologies, and she is an active member of IEEE.



HASAN MAHMUD SOZIB received the B.Eng. degree in Electrical and Electronic Engineering from Ahsanullah University of Science and Technology (AUST), Dhaka, Bangladesh, and the MBA degree in Human Resource Management from North South University, Dhaka. He is currently serving as a Senior Executive in Human Resources at Al-Mostafa Group, where he is engaged in performance management, recruitment, and employee development. He previously worked at Mensa DigiWorld and the U.S. Embassy, Dhaka, and also co-founded BizCord Digital. His professional and research interests span artificial intelligence, machine learning, and Internet-of-Things (IoT) applications, with contributions to peer-reviewed publications in these fields.



HALIMA AHMED is pursuing a Bachelor of Science in Public Health at California State University, Los Angeles. She previously completed an Associate of Arts in Liberal Arts: Natural Applied Sciences, an Associate of Science in Pre-Allied Health, and an AS Transfer in Public Health Science from Los Angeles City College in 2025. Her research interests include environmental health, food safety and public health policy, community health equity, healthcare management, and public health data analysis. She is dedicated to advancing sustainable health solutions and promoting equity through research and practice.



AFIA FAIROOZ TASNIM is pursuing a Master of Public Health with a concentration in Community Health at California State University, Long Beach, where she champions health equity and culturally responsive care. She holds a B.Sc. in Public Health (minor in Mathematics) from the Asian University for Women. Afia has engaged in impactful public health work, including analyzing service trends using Tableau and Power BI during her internship at the Community Action Partnership of Orange County, designing an intervention-focused program proposal to address perinatal depression, and mapping active commuting behaviors using ArcGIS. Her research has also explored the relationship between adolescent screen time and mental health, as well as structured interventions tackling opioid use among pregnant women and physical inactivity in adults. Through these experiences, she has developed strong competencies in data analysis (SPSS, GIS), program planning, research design, and community-driven health promotion. She is committed to advancing inclusive, evidence-based, and scalable public health interventions to improve outcomes for underserved populations.



ADIB HOSSAIN is pursuing a Master of Science in Business Analytics at Trine University, where he maintains a strong academic record with a GPA of 3.75/4. He previously earned his B.Sc. in Apparel Manufacturing & Technology from BGMEA University of Fashion & Technology (BUFT) with distinction. Professionally, he works as a Junior Supply Chain Analyst at NYX LLC and has prior experience as a Junior Production Merchandiser at Padma Textile Ltd. He is also a Certified Supply Chain Analyst (CSCA™) from the International Supply Chain Education Alliance, USA. His research and professional interests include business analytics, supply chain management, and digital solutions for operational efficiency.



FAHAD AHMED is a motivated Quality Engineer currently pursuing a Master of Science in Engineering Management at Trine University, Indiana, with a strong academic record (CGPA 3.87/4.00). He previously earned a Bachelor of Science in Textile Engineering from BGMEA University of Fashion Technology, Bangladesh, including an exchange program at DKTE Textile Engineering Institute, India. His professional experience includes roles as Engineering Intern at NYX LLC in Michigan and Assistant Executive in Quality Control at SGS Bangladesh. Certified in Lean Six Sigma and Project Management, Fahad is also a student member of the American Society for Quality (ASQ) and the Project Management Institute (PMI). His research and professional interests focus on quality assurance, process improvement, and operational efficiency.



ANSEENA ANEES SABEENA is an internationally experienced business professional currently pursuing a Doctor of Business Administration at Westcliff University, Irvine, California. She holds a Global MBA from De Montfort University, England, and a Bachelor of Business Administration from Calicut University, India. Her professional experience spans marketing management, education consulting, healthcare training, and network marketing, with proven expertise in international marketing, customer relationship management, and leadership. She has also contributed to global business projects focusing on sustainability and organizational growth. Fluent in English, Malayalam, and Tamil, she is dedicated to advancing research and practice in strategic leadership, marketing, and business development.



MD IMRAN SARKAR is currently pursuing a Master of Science in Information Technology at Westcliff University. He holds a Bachelor of Science in Computer and Communication Engineering, which provides him with a strong foundation in technology and engineering principles. With a keen interest in advanced computing, networking, and information systems, Md Imran Sarkar is dedicated to exploring innovative solutions in the field of information technology. His academic journey reflects a commitment to excellence and a passion for contributing to technological advancements that address real-world challenges.



TIMOTEI ISTVÁN ERDEI is an Assistant lecturer, Mechatronical Engineer (BSc), Mechatronical Engineer (MSc), Lean Engineer, Quality Control Engineer, Doctoral School of Informatics, PhD studies in the field of industrial robotics. Several years of experience as a departmental engineer, research and development, and application of innovative solutions (innovator) at the University of Debrecen. His main research areas are Cyber-Physical Systems, Augmented Reality,

Virtual Reality, AI/artificial intelligence, Self-driving autonomous vehicles (AGV), Image processing, Industrial robotics, 3D printer / CNC design and implementation, industrial robot simulations, production line optimization, PLC programming, PCB design, automated industrial technical systems, ROS Os - Linux operating systems.